

AUTO-PRUNE: Automated DNN Pruning and Mapping for ReRAM-Based Accelerator

Siling Yang^{*}, Weijian Chen^{*}, Xuechen Zhang[#], Shuibing He^{*},
Yanlong Yin^{\$}, Xian-He Sun⁺

*



浙江大学
Zhejiang University

#



\$



ZHEJIANG LAB
之江实验室

+








ILLINOIS TECH

Accelerating the DNN

**Deep neural network (DNN)
is popular in various fields.**

➤ DNN Hardware accelerator

- GPU
- TPU
- ASIC
- Novel architectures and emerging devices



...

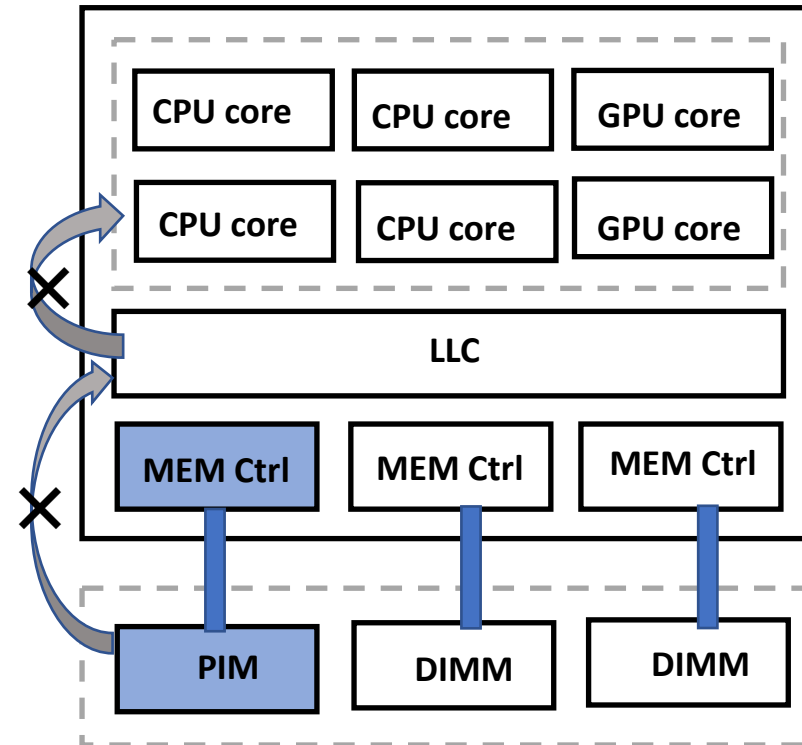
Von Neumann Architecture vs. Processing-in-Memory



Energy Wall


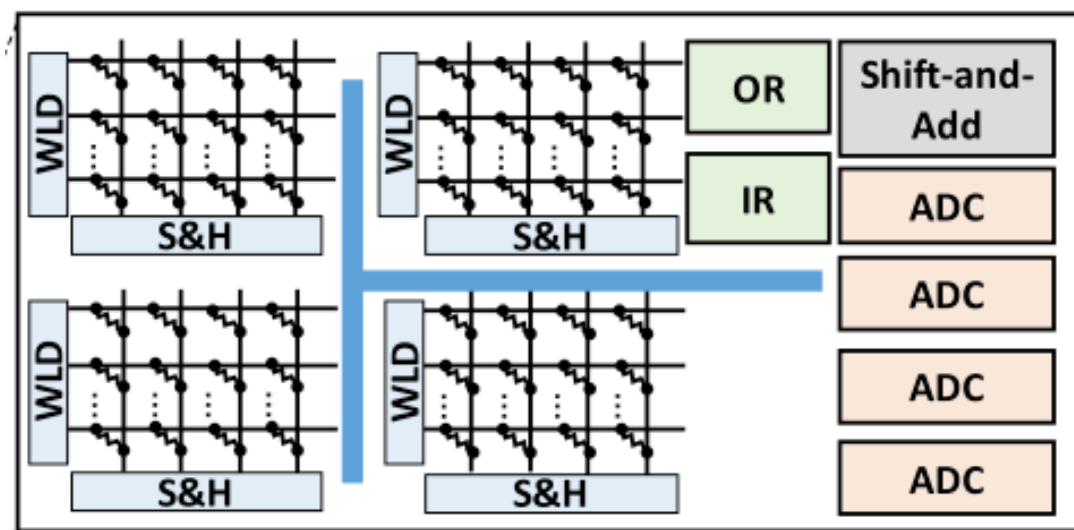
| Operation | Energy(pJ) |
|--------------------|------------|
| 16b Add | 0.05 |
| 32b Add | 0.1 |
| 16b FP Add | 0.4 |
| 32b FP Add | 0.9 |
| 32b Mult | 3.1 |
| 16b FP Mult | 1.1 |
| 32b FP Mult | 3.7 |
| 32b SRAM Read(8KB) | 5 |
| 32b DRAM Read | 640 |

Processing in Memory




PIM and emerging devices can alleviate the energy wall.

ReRAM-based Accelerator



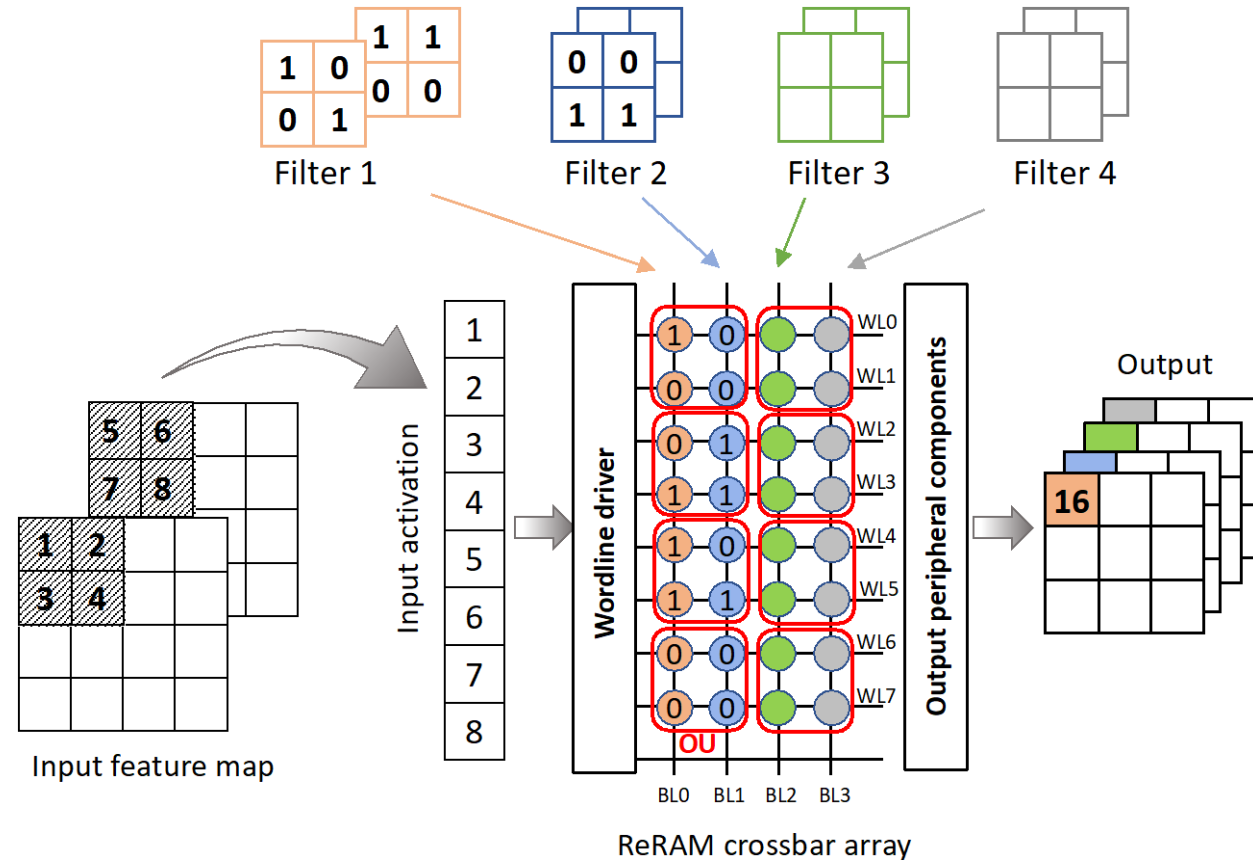
(a) Multiply and Accumulation



(b) Matrix-Vector Multiplication

ReRAM-based DNN accelerator architecture.[SRE-19]

Mapping Filter Weights of DNNs in ReRAM-based Accelerators



➤ Operation Unit (OU)

Illustration of mapping filter weights to a crossbar array used in the architecture of ReRAM-based accelerators.

Filter weight matrices of DNN models are sparse.

Related Work & Motivation

| Pruning techqiuie | Method | Hardware customization | Pattern for pruning | Use OU in data-path |
|--------------------------|------------|------------------------|------------------------------|---------------------|
| LSR[ASPDAC19] | heuristics | ✗ | Unimportant weight groups | ✗ |
| SRE[ISCA20] | heuristics | ✗ | All-zero row/column vectors | ✓ |
| PIM-Prune[DAC20] | heuristics | ✗ | Unimportant rows and columns | ✗ |
| Pattern pruning[arxiv20] | heuristics | ✗ | Patterns | ✓ |

1. They use heuristics to prune the weights, leading to **suboptimal pruning policies**.
2. They mostly focus on improving compression ratio, thus **may not meet accuracy constraints**.
3. They **ignore direct feedback of hardware**, e.g., the number of occupied crossbars or energy consumption.

Objectives of Our Work

Make a
global optimal
pruning policy

+

Make a pruning
and mapping
policy tailored for
different hardware

+

Avoid the
dislocation
problem



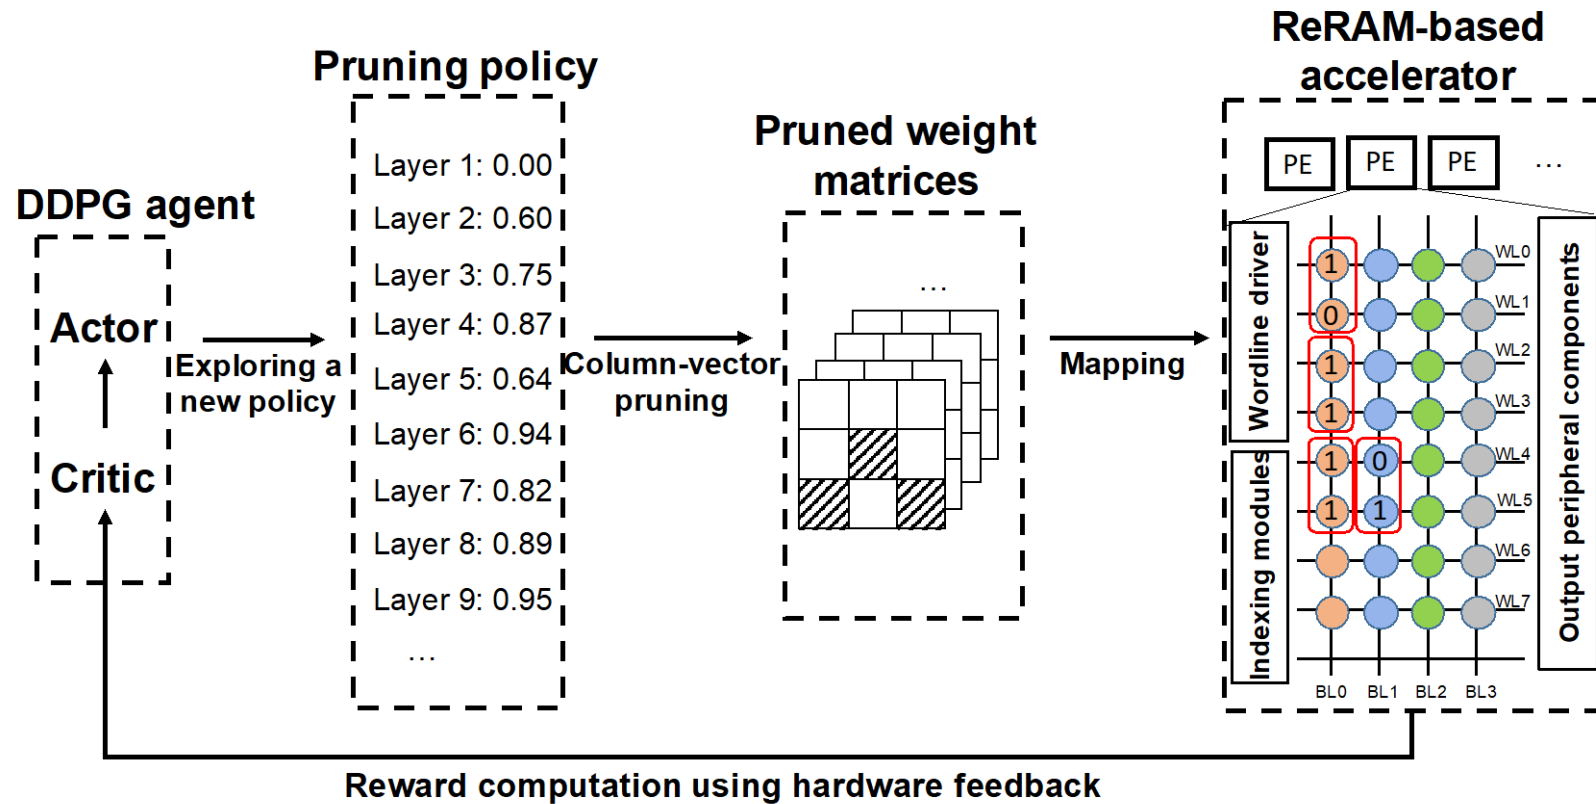
AUTO-PRUNE

Design of AUTO-PRUNE

Main Design

- DDPG Algorithm for ReRAM-based Accelerator
- Column-Vector Based Pruning and OU Formation
- Data-Path Design

Overview of AUTO-PRUNE



1. DDPG Algorithm for ReRAM-based Accelerator

- State Space: identify a layer with its characteristics

$$(k, t, inc, outc, ks, h, w, s, xb[k], xb_{saved}[k], xb_{rest}[k], a_{k-1})$$

- Action Space: pruning rate for a specified layer

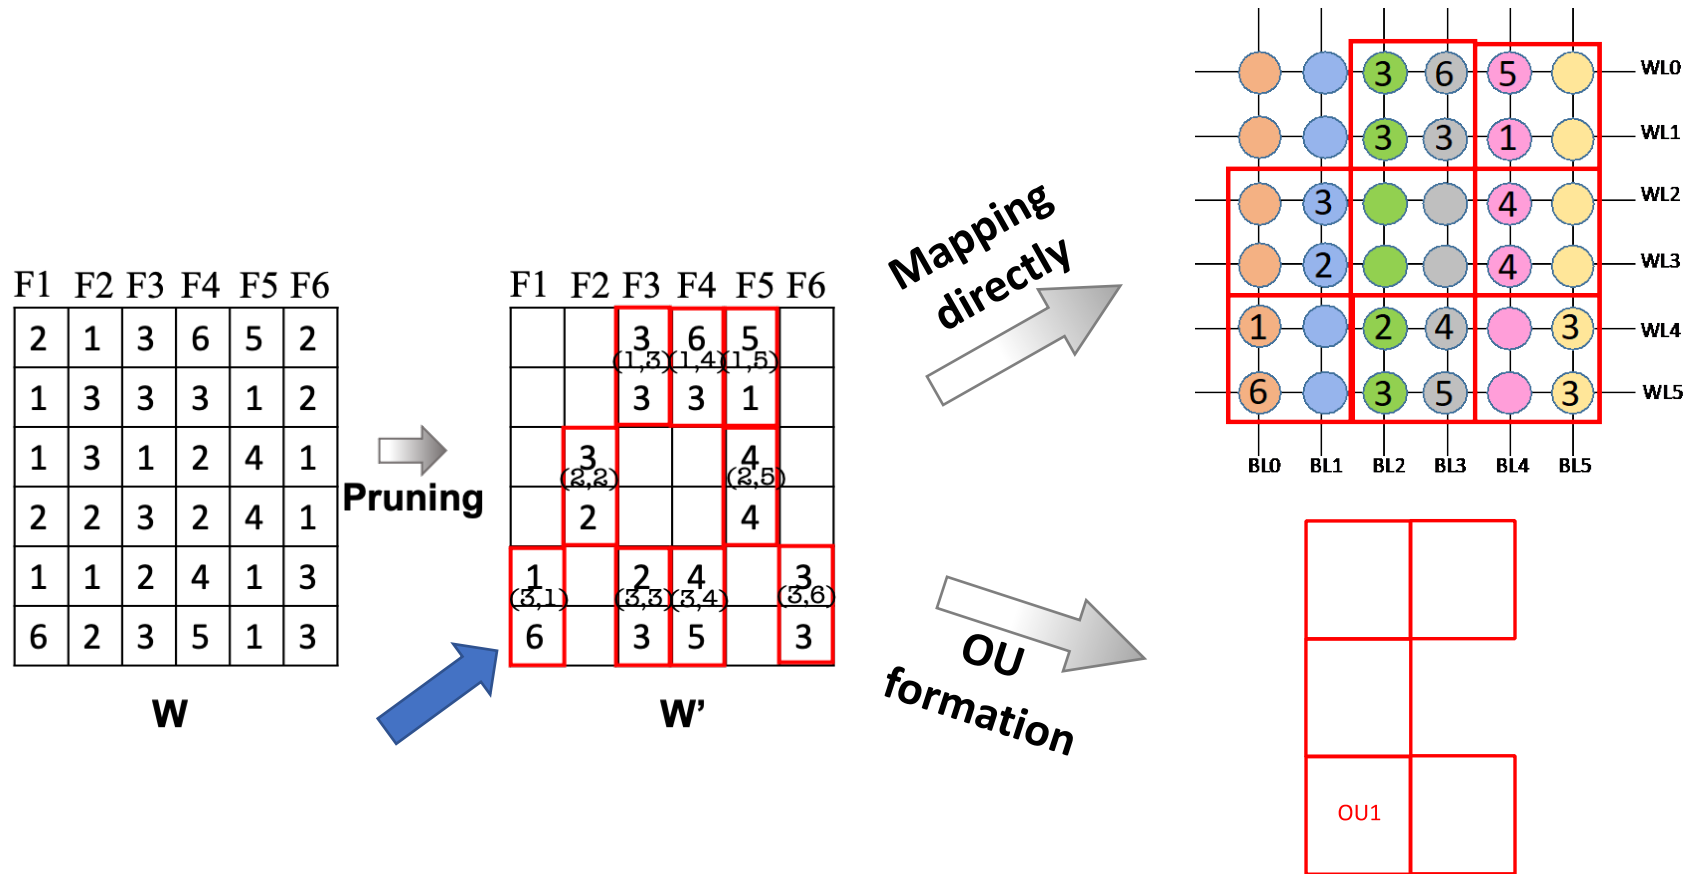
- $a_k \in (0, 1]$

- Reward Function: related with compression rate and accuracy

$$Reward = (1 - \frac{1}{rate_{compression}^{xb}})^{\alpha} \times acc_{reram}$$

| Symbol | Meaning |
|-----------------|--|
| k | layer index |
| t | layer type: CONV:1 ; FC: 0 |
| inc | number of channels in the input feature map |
| $outc$ | number of channels produced by the convolution |
| ks | number of elements of a convolving kernel |
| h | height of the input feature maps |
| w | width of the input feature maps |
| s | stride of the convolution |
| $xb[k]$ | number of crossbars required for mapping layer k |
| $xb_{saved}[k]$ | accumulated number of the crossbar saved from the first layer to layer $k - 1$ |
| $xb_{rest}[k]$ | number of crossbars required from layer $k + 1$ to the last layer |
| $size_{xb}$ | length of the crossbar size |
| acc_{reram} | accuracy reported by the ReRAM-based accelerator simulator |
| a_{k-1} | action from the last time step |

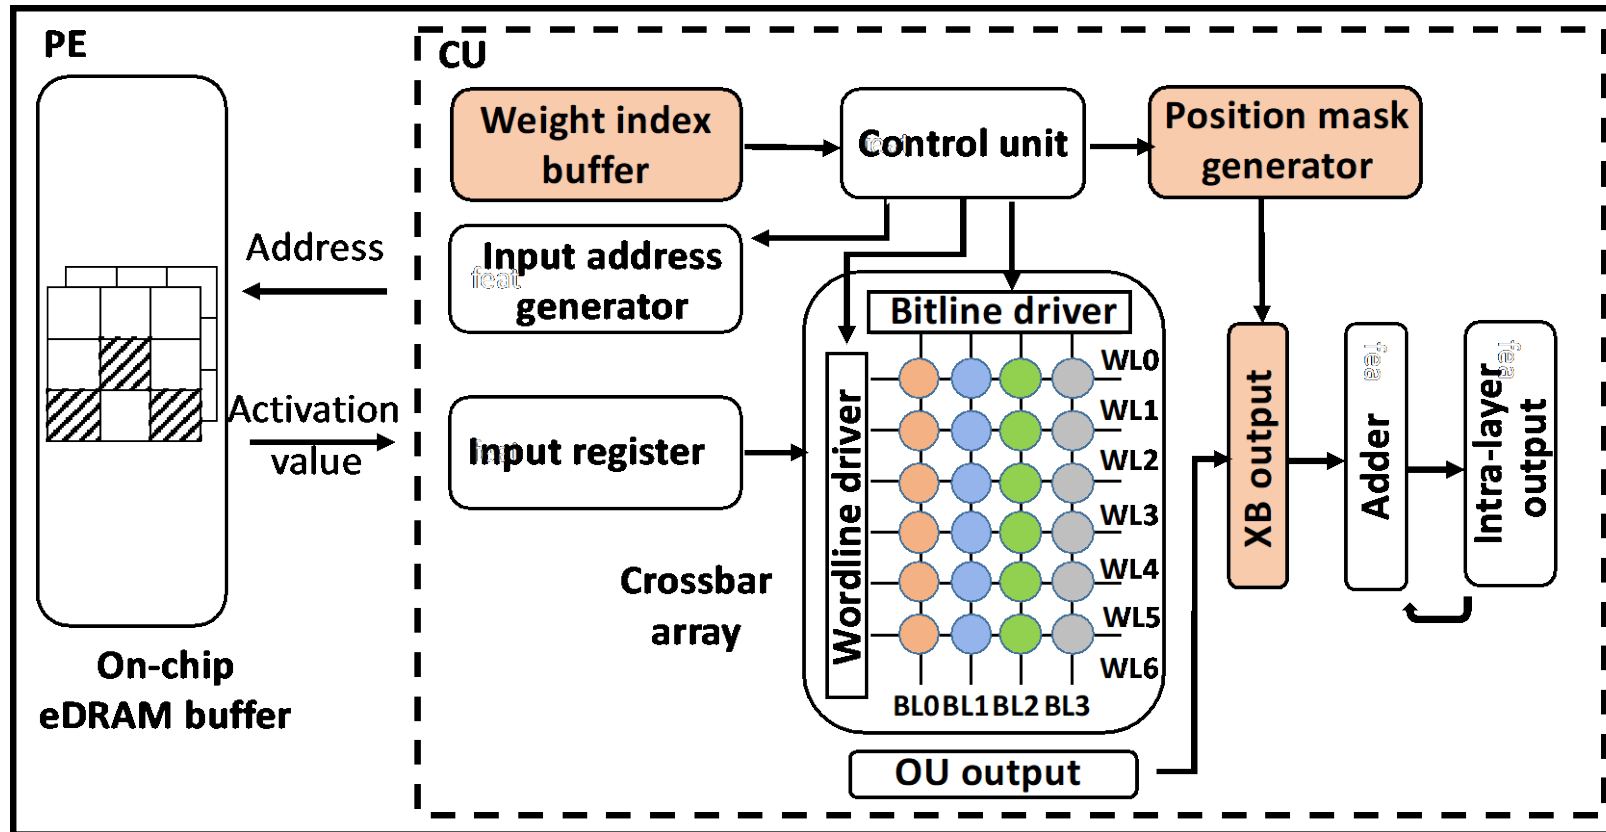
2. Column-Vector Based Pruning and OU Formation



➤ OU List

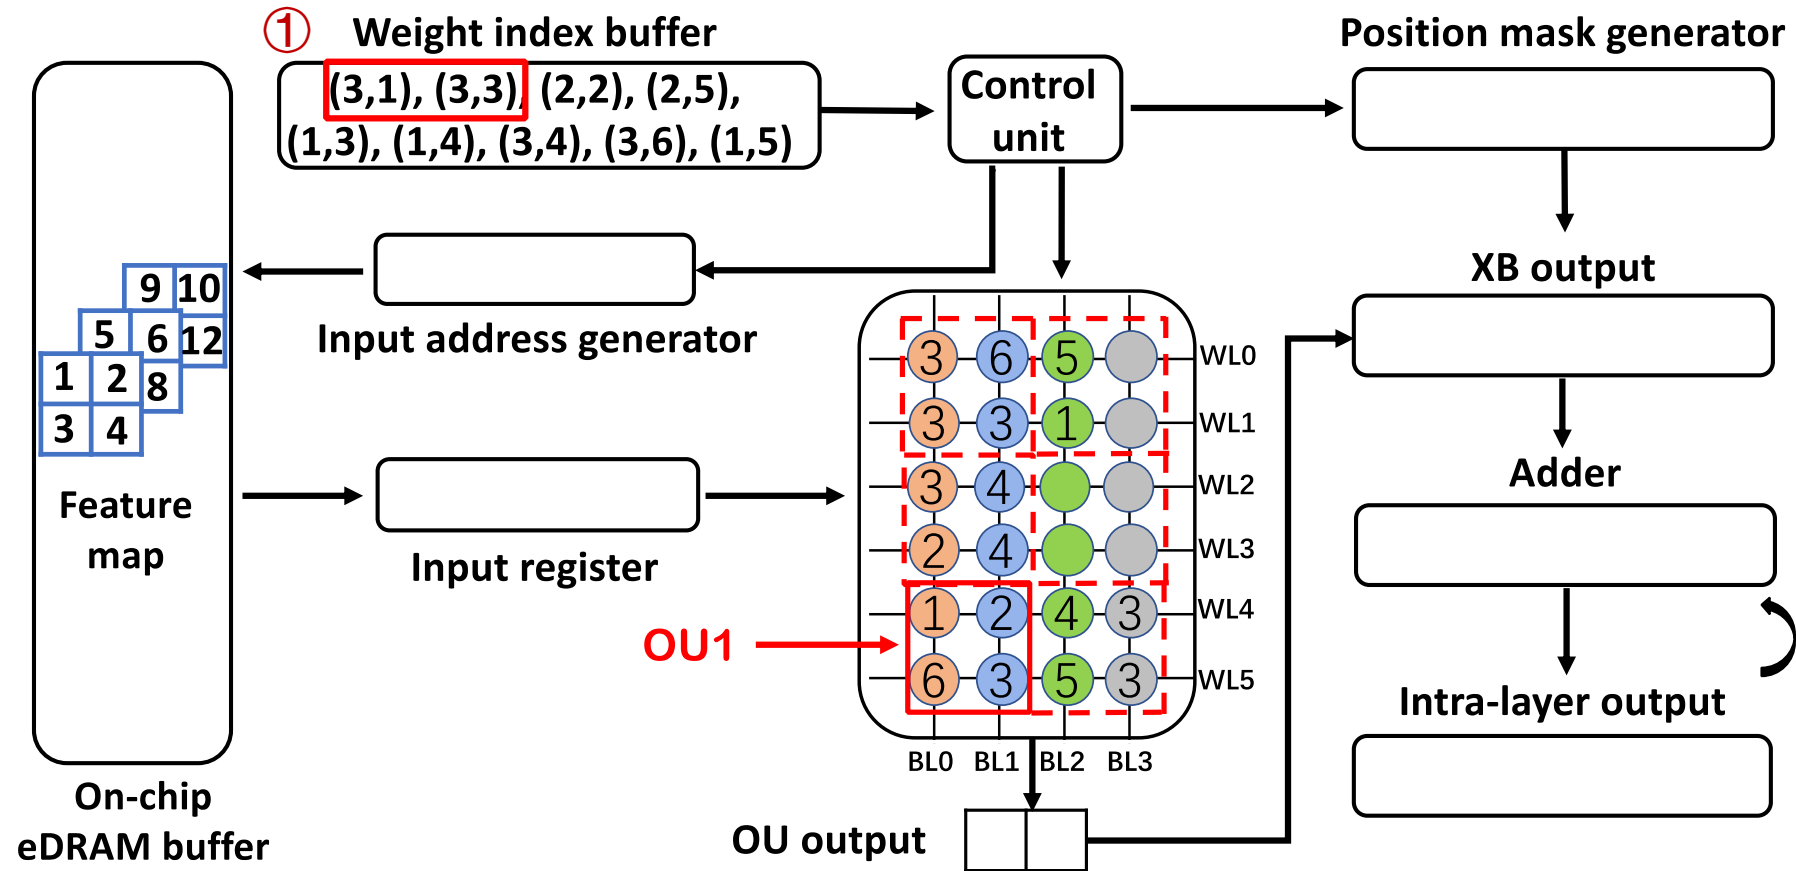
$\{(3,1), (3,3), (2,2), (2,5), (1,3), (1,4), (3,4), (3,6), (1,5)\}$

3. Data-Path Design




An overview of the data-path for Auto-prune.

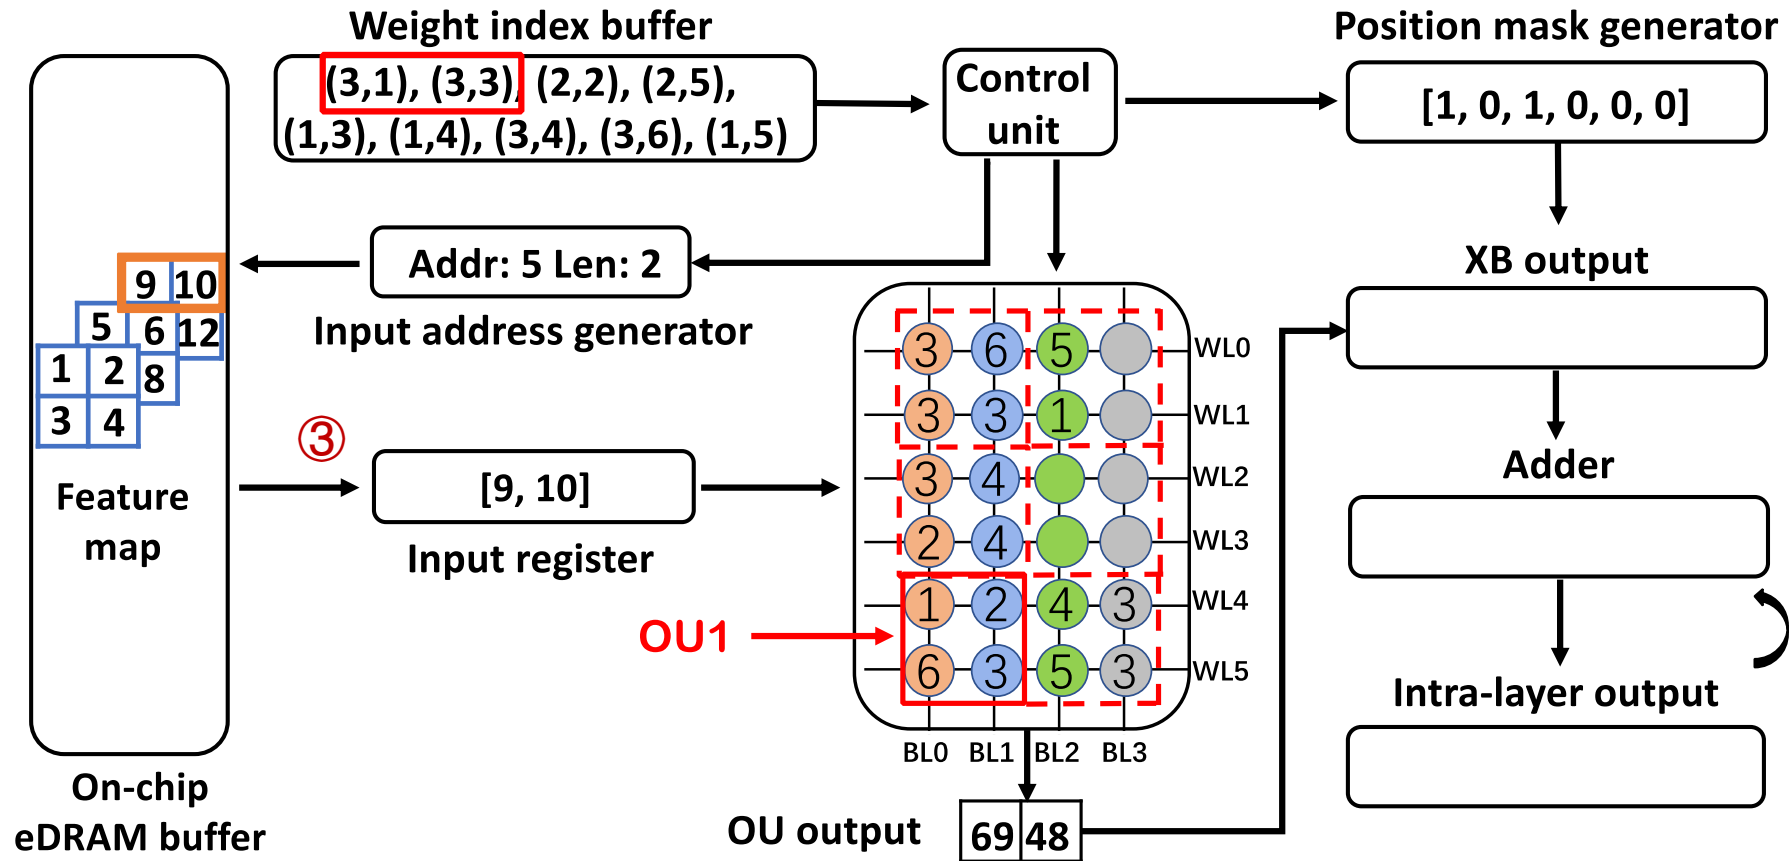
An example: conv operation in OU1



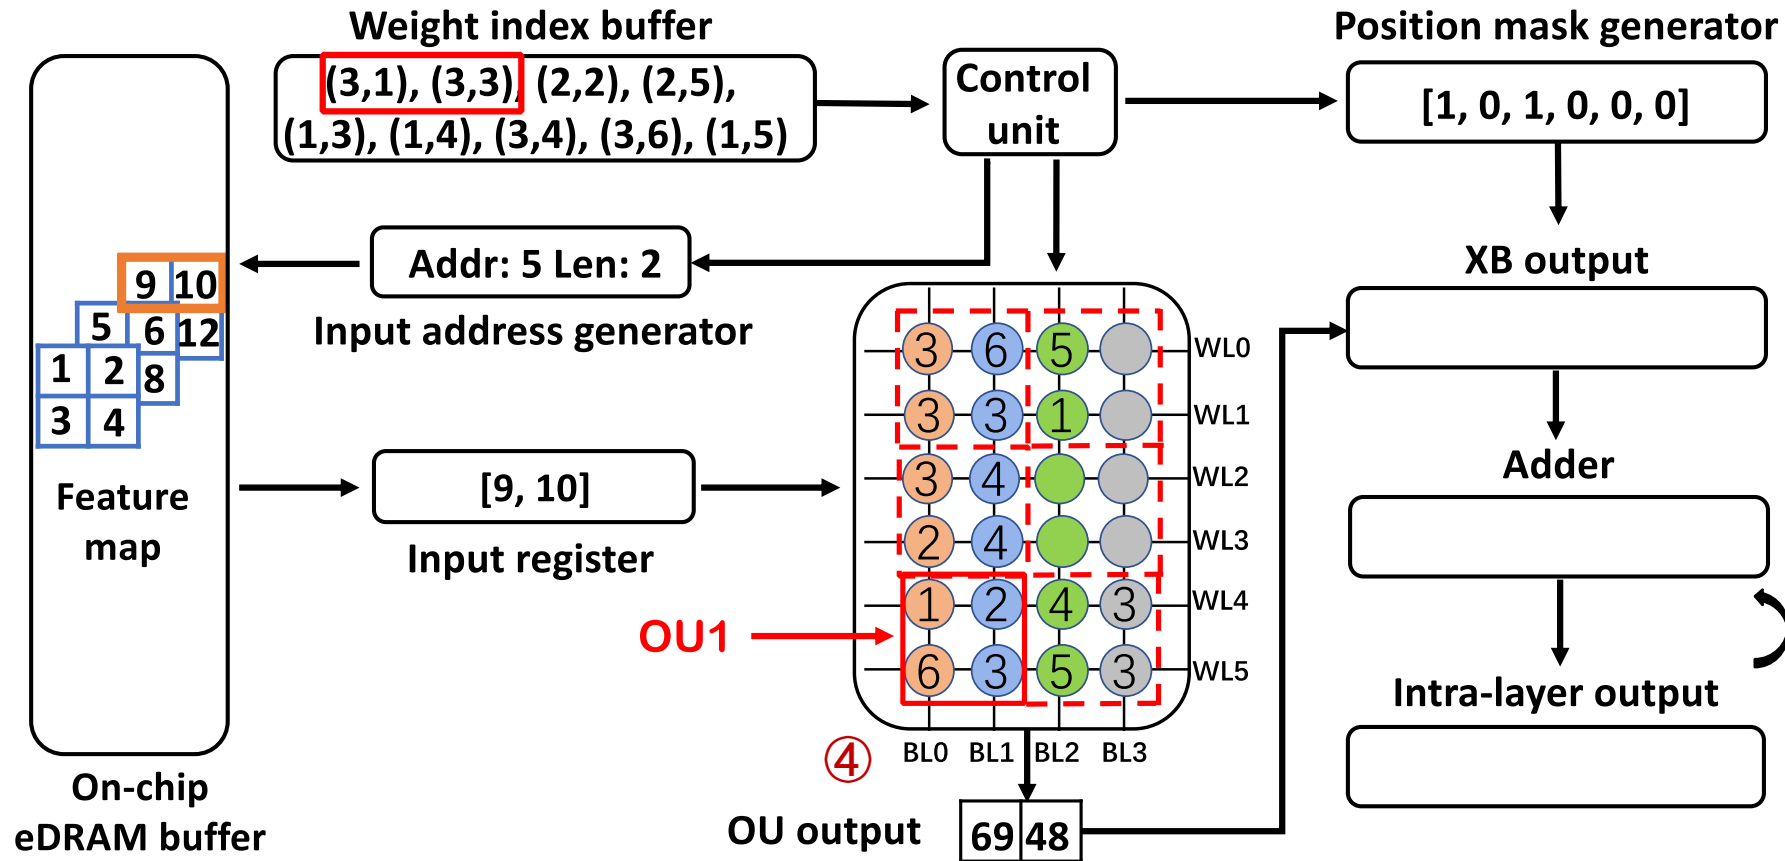
An example: conv operation in OU1



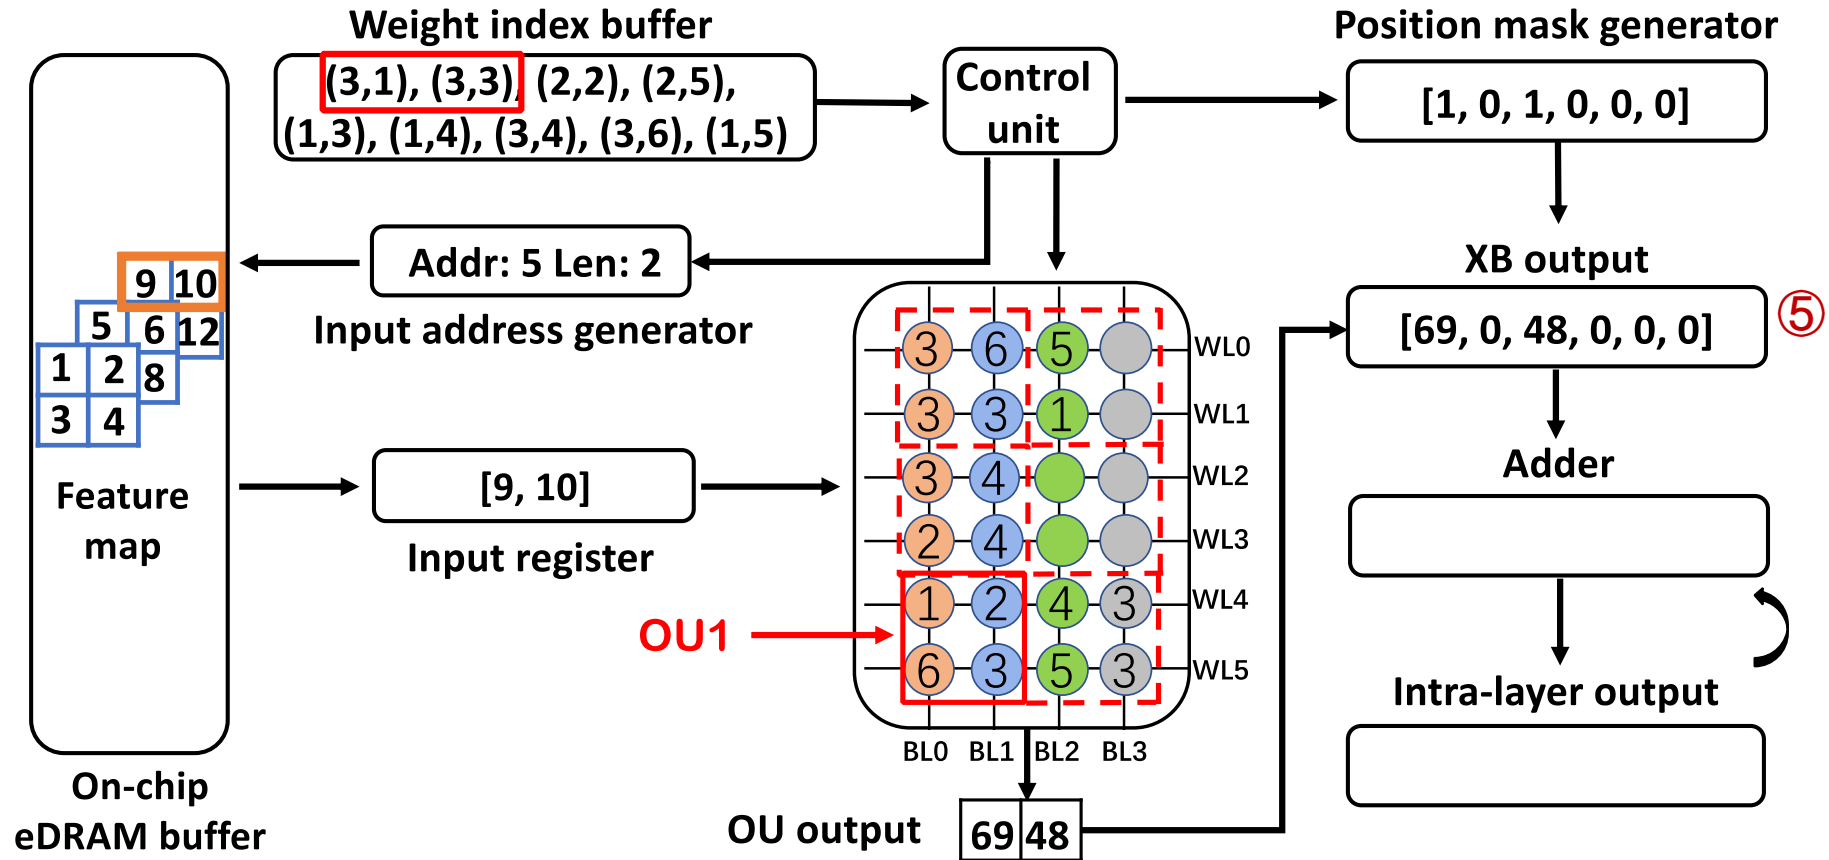
An example: conv operation in OU1



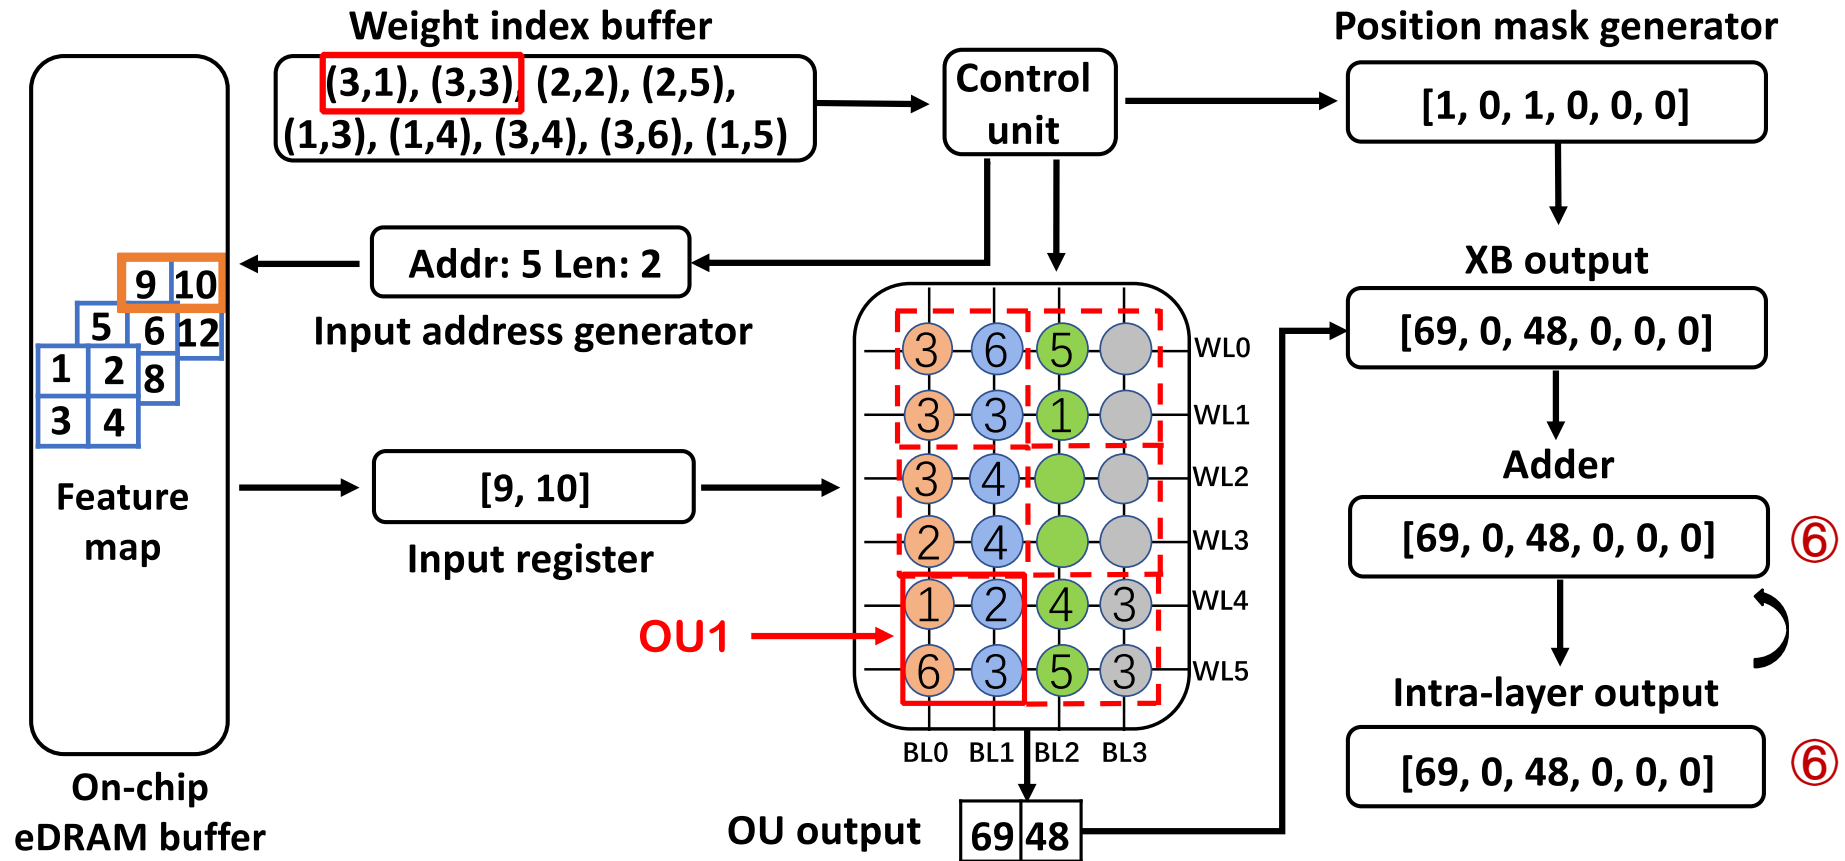
An example: conv operation in OU1



An example: conv operation in OU1



An example: conv operation in OU1



Experimental Setting

➤ simulator: MNSIM-2.0

- Crossbar size: 128 x 128
- bit-per-cell: 1 bits
- OU size: 32 x 32

➤ Counterparts

- Naïve
- PIM-Prune[DAC-20]
- Pattern-Prune[Arxiv-20]

➤ Workloads and datasets

- NN: AlexNet, VGG16, Plain20
- Dataset: CIFAR10, MNIST

➤ Metrics

- Compression rate
- Energy & area efficiency

➤ Discussion


- Sensitivity & Overhead

Compression rate

| Network | Method | CR on XBs | Acc5 | Acc Drop |
|---------|---------------|-----------|--------|----------|
| AlexNet | Naïve | 1 | 99.36% | - |
| | PIM-Prune | 4.3 | 98.81% | 0.55% |
| | Pattern-Prune | 1.1 | 96.48% | 2.88% |
| | Auto-Prune | 14.3 | 99.10% | 0.26% |
| VGG16 | Naïve | 1 | 99.29% | - |
| | PIM-Prune | 6.1 | 98.62% | 0.67% |
| | Pattern-Prune | 2.6 | 98.43% | 0.86% |
| | Auto-Prune | 11.9 | 98.62% | 0.67% |
| Plain20 | Naïve | 1 | 98.14% | - |
| | PIM-Prune | 7.3 | 98.19% | -0.05% |
| | Pattern-Prune | 1.2 | 98.24% | -0.10% |
| | Auto-Prune | 10.3 | 98.29% | -0.15% |

the same or higher accuracy, compression rate up to:
3.3X PIM-Prune
13X Pattern-Prune


Energy efficiency & area efficiency



(a) CIFAR10.

the result of energy efficiency on CIFAR10

12.2 Pattern-Prune
2.3 PIM-Prune



(a) CIFAR10.

the result of area efficiency on CIFAR10

3.1X Pattern-Prune
9.6X PIM-Prune

Sensitivity study

- granularity of column-vector
-




Figure 10: Compression rate, area efficiency and energy efficiency for AlexNet with various granularities of column-vectors.

The smaller granularity of column-vector, the higher compression rate.

Index overhead




Figure 12: The storage overhead of Weight Index Buffer for different networks on CIFAR10 and MNIST respectively.

The index overhead is ignorable.

Conclusions

- AUTO-PRUNE is a hardware-aware automated DNN pruning and mapping framework for ReRAM-based accelerators. It leverages RL to automatically determine a global optimum pruning policy, considering the direct hardware feedback.
- We propose a new data-path to correctly index and feed input to matrix-vector computation.
- AUTO-PRUNE achieves up to 3.3X compression rate, 3.1X area efficiency, and 3.3X energy efficiency compared to PIM-Prune while maintaining a similar or even higher accuracy.

Thanks for your attention!

Siling Yang@ZJU
slingzjunet@zju.edu.cn



浙江大学
Zhejiang University