



# Efficient Maximal Biclique Enumeration on GPUs

Zhe Pan, Shuibing He, Xu Li, Xuechen Zhang\*, Rui Wang, Gang Chen

Zhejiang University, \*Washington State University Vancouver



# Outline

## ➤ Introduction

- Problem definition
- MBE on CPUs
- Related work comparison

## ➤ Challenges of MBE on GPUs

- Large memory requirement
- Massive thread divergence
- Load imbalance

## ➤ GMBE : the **first** highly-efficient GPU solution for the MBE problem

- Stack-based iteration with **node reuse**
- Pruning using **local neighborhood sizes**
- **Load-aware** task scheduling

## ➤ Evaluation

# Introduction

- Problem definition
- MBE on CPUs
- Related work comparison

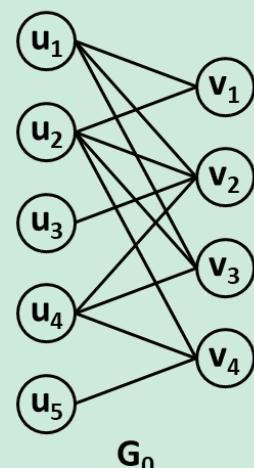
# Introduction : Problem Definition

## ➤ Preliminaries

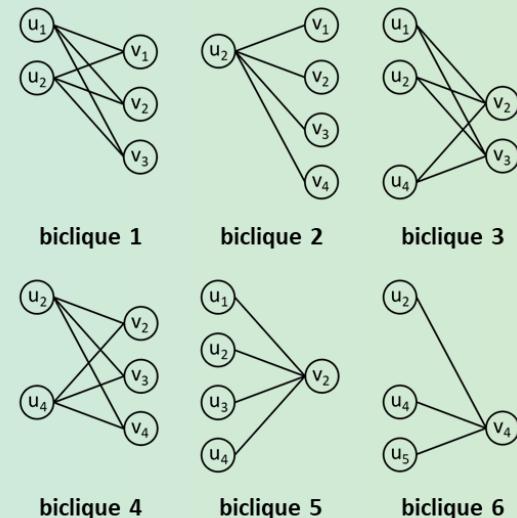
- **Bipartite graph**  $G(U, V, E)$  : A graph structure contains two disjoint vertex sets  $U, V$  and an edge set  $E. E \subseteq U \times V$ .
- **Biclique** : A complete bipartite graph in which every vertex is connected to every vertex in the opposite subset.
- **Maximal biclique** : a biclique that can not be further enlarged to form a large biclique.

## ➤ Problem definition

- Maximal biclique enumeration (MBE) aims to find all maximal bicliques in  $G$ .



A bipartite graph  $G_0$  containing 6 maximal bicliques.



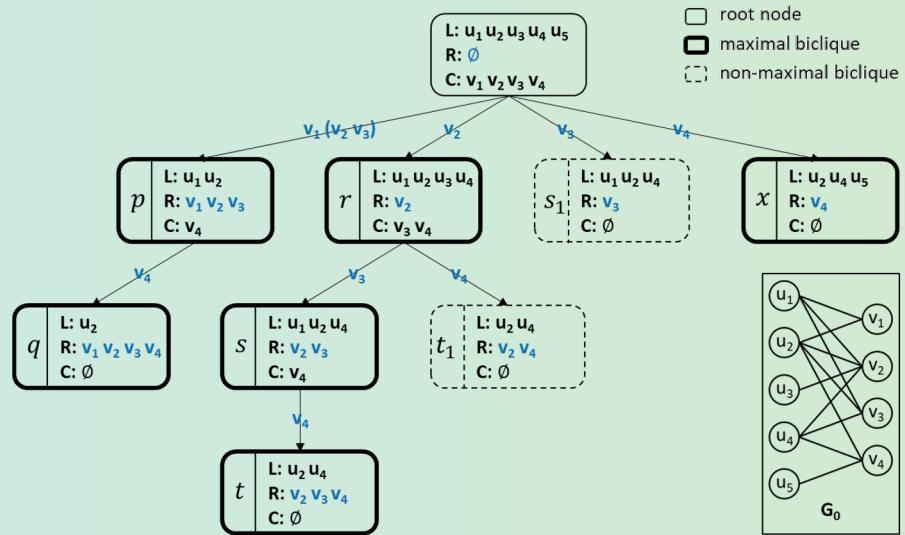
# Introduction : MBE on CPUs

## ➤ Set enumeration tree for MBE

- Each tree node is a 3-tuple  $(L, R, C)$ .  $(L, R)$  is the corresponding biclique and  $C$  stores candidate vertices for expanding  $R$ .

## ➤ Baseline solution

- Step 1 : Utilize a set enumeration tree to generate **the powerset of  $V$** .
- Step 2 : Expand each subset of the powerset of  $V$  to a biclique  $(L, R)$  and enumerate maximal ones.



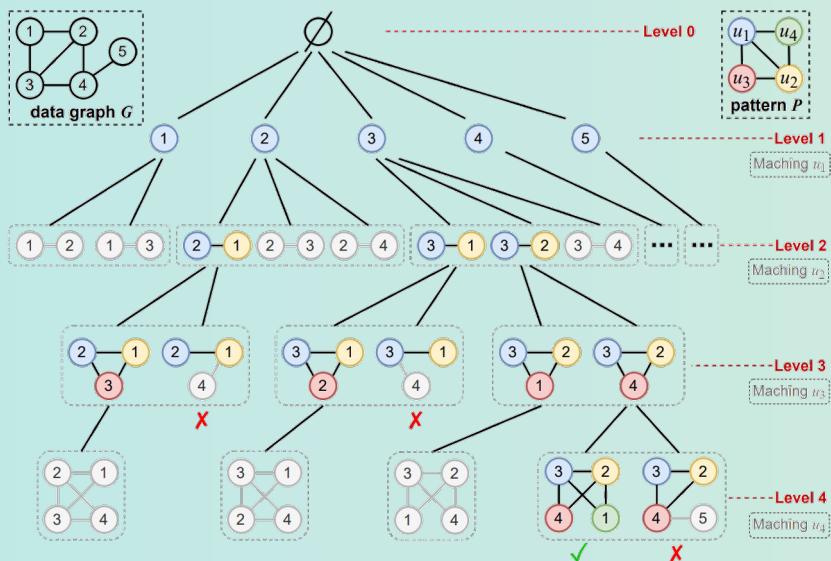
# Introduction : MBE on CPUs

- **Recent optimizations**
  - Vertex ordering [1, 2, 5]
  - Candidates pruning using pivots [1, 2]
  - Parallelization on multicore CPUs [3] or distributed architectures [4]

**Existing solutions for MBE are insufficient because their performance speedup is constrained by the limited parallelism of CPUs.**

- [1] Lu Chen, Chengfei Liu, Rui Zhou, Jiajie Xu, and Jianxin Li. 2022. Efficient Maximal Biclique Enumeration for Large Sparse Bipartite Graphs. VLDB 2022. 1559-1571.
- [2] Aman Abidi, Rui Zhou, Lu Chen, and Chengfei Liu. Pivot-Based Maximal Biclique Enumeration. IJCAI 2020. 3558–3564.
- [3] Apurba Das and Srikanta Tirthapura. 2019. Shared-Memory Parallel Maximal Biclique Enumeration. HiPC 2019.
- [4] Arko Provo Mukherjee and Srikanta Tirthapura. Enumerating Maximal Bicliques from a Large Graph Using MapReduce. IEEE Trans. Serv. Comput. 10, 5 (2017), 771–784.
- [5] Yun Zhang, Charles A. Phillips, Gary L. Rogers, Erich J. Baker, Elissa J. Chesler, and Michael A. Langston. BMC bioinformatics 15, 1 (2014), 110.

# Introduction : Related Work Comparison



An enumeration tree for mining pattern  $P$  in data graph  $G$ .

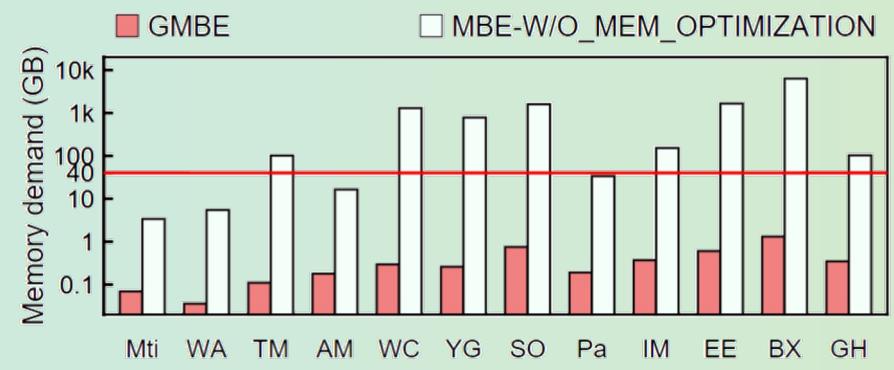
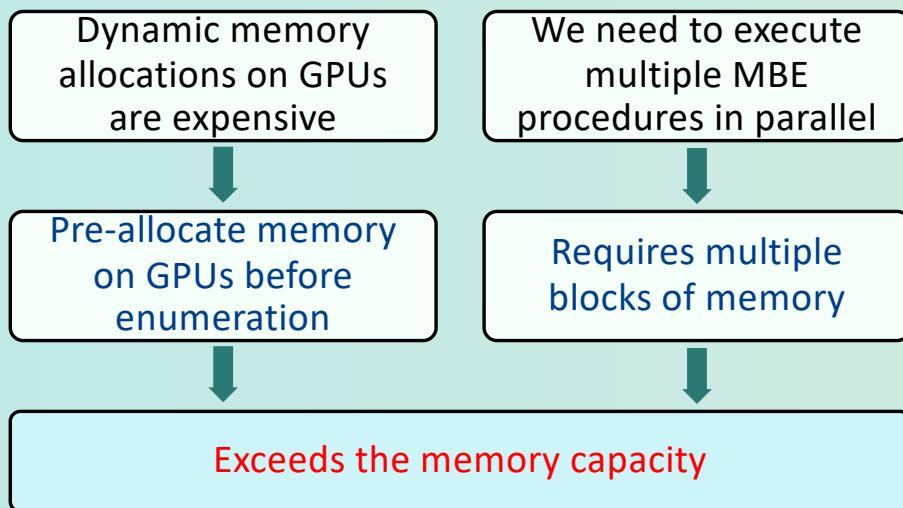
Problem	MBE	Graph pattern mining [1]
Vertex count in enumerated subgraphs	Unfixed number of vertices, can be large.	Fixed number of vertices equivalent to pattern size $ P $ , typically small.
Enumeration tree height	Unfixed and can be up to $d_{max}(V)$ .	Fixed and equal to $ P $ .
Conclusion	(1) MBE requires <b>significantly more memory</b> than GPM to actively maintain up to $d_{max}(V)$ tree nodes for backtracking. (2) MBE generates <b>more severe imbalanced workloads</b> than GPM due to the variation in height among its enumeration trees.	

[1] Xuhao Chen and Arvind. Efficient and Scalable Graph Pattern Mining on GPUs. OSDI 2022. 857–877.

# Challenges

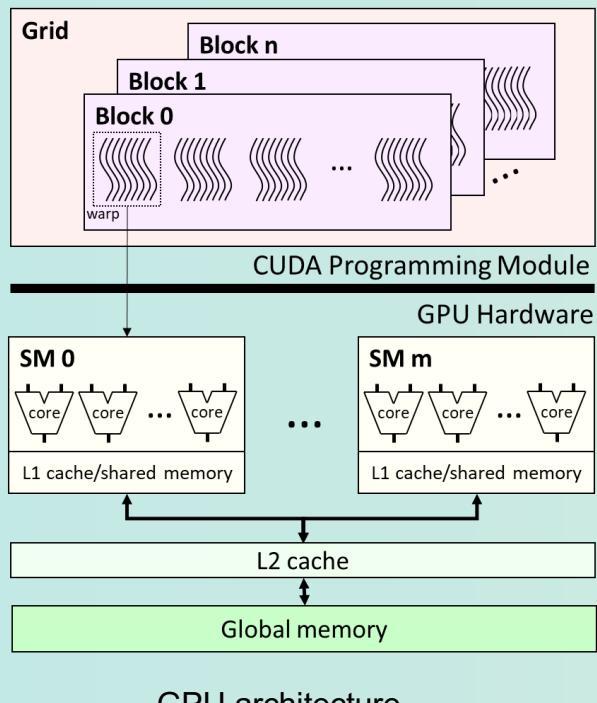
- Large memory requirement
- Massive thread divergence
- Load imbalance

# Challenge 1 : Large Memory Requirement



Directly parallelizing existing MBE procedures on an A100 GPU will exceed the memory capacity on multiple datasets.

# Challenge 2 : Massive Thread Divergence



CS 0

```
if((threadIdx.x & 1) == 1){  
    code A;  
    if((threadIdx.x & 2) == 2){  
        code B;  
    } else{  
        code C;  
    }  
}else{  
    code D;  
    if((threadIdx.x & 2) == 2){  
        code E;  
    } else{  
        code F;  
    }  
}
```

CS 1

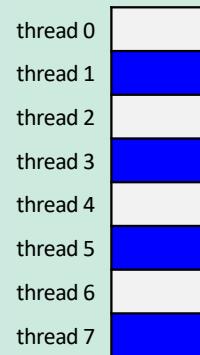
```
code A;  
code A;
```

CS 0 and CS 1 are GPU code segments where threads with different routines execute **2 sets of codes** each.

# Challenge 2 : Massive Thread Divergence

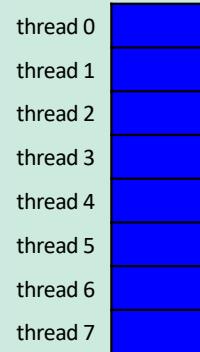
CS 0

```
if((threadIdx.x & 1) == 1){  
    code A;  
    if((threadIdx.x & 2) == 2){  
        code B;  
    } else{  
        code C;  
    }  
} else{  
    code D;  
    if((threadIdx.x & 2) == 2){  
        code E;  
    } else{  
        code F;  
    }  
}
```



CS 1

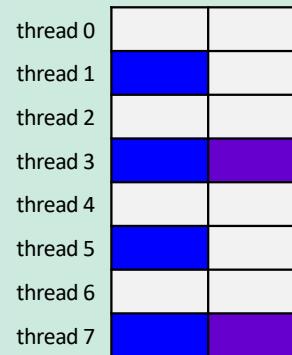
```
code A;  
code A;
```



# Challenge 2 : Massive Thread Divergence

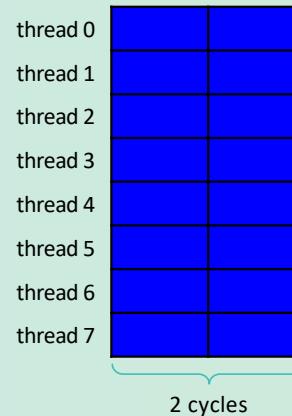
CS 0

```
if((threadIdx.x & 1) == 1){  
    code A;  
    if((threadIdx.x & 2) == 2){  
        code B;  
    } else{  
        code C;  
    }  
} else{  
    code D;  
    if((threadIdx.x & 2) == 2){  
        code E;  
    } else{  
        code F;  
    }  
}
```



CS 1

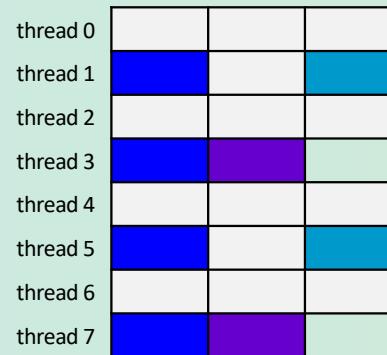
```
code A;  
code A;
```



# Challenge 2 : Massive Thread Divergence

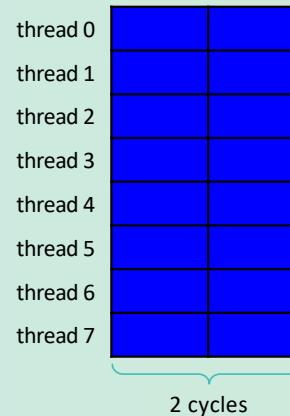
CS 0

```
if((threadIdx.x & 1) == 1){  
    code A;  
    if((threadIdx.x & 2) == 2){  
        code B;  
    } else{  
        code C;  
    }  
} else{  
    code D;  
    if((threadIdx.x & 2) == 2){  
        code E;  
    } else{  
        code F;  
    }  
}
```



CS 1

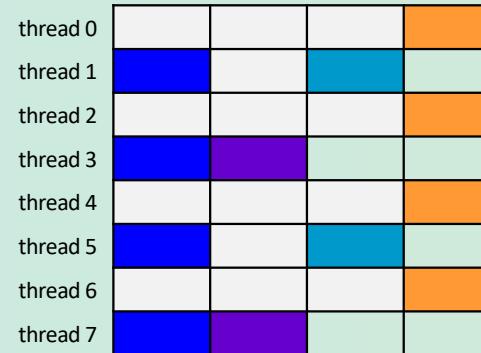
```
code A;  
code A;
```



# Challenge 2 : Massive Thread Divergence

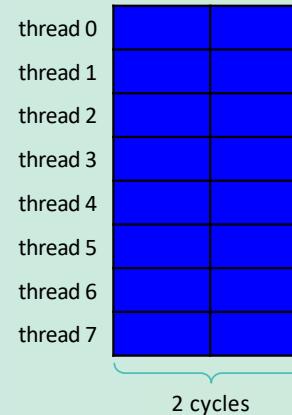
CS 0

```
if((threadIdx.x & 1) == 1){  
    code A;  
    if((threadIdx.x & 2) == 2){  
        code B;  
    } else{  
        code C;  
    }  
} else{  
    code D;  
    if((threadIdx.x & 2) == 2){  
        code E;  
    } else{  
        code F;  
    }  
}
```



CS 1

```
code A;  
code A;
```



# Challenge 2 : Massive Thread Divergence

CS 0

```
if((threadIdx.x & 1) == 1){  
    code A;  
    if((threadIdx.x & 2) == 2){  
        code B;  
    } else{  
        code C;  
    }  
} else{  
    code D;  
    if((threadIdx.x & 2) == 2){  
        code E;  
    } else{  
        code F;  
    }  
}
```

thread 0					
thread 1	Blue		Cyan	Light Green	Dark Green
thread 2				Orange	Dark Green
thread 3	Blue	Purple	Light Green	Light Green	
thread 4				Orange	
thread 5	Blue		Cyan	Light Green	Light Green
thread 6				Orange	Dark Green
thread 7	Blue	Purple			

CS 1

```
code A;  
code A;
```

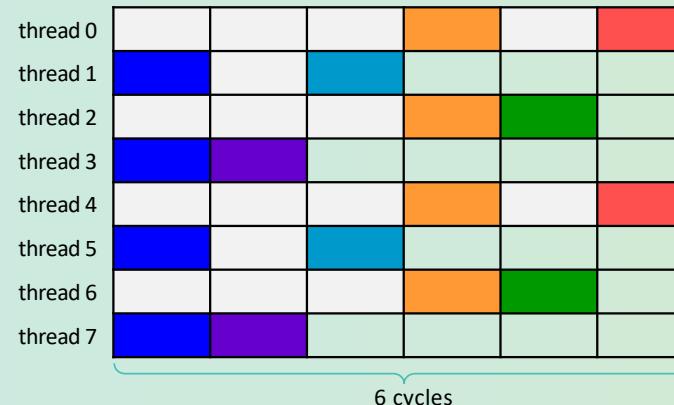
thread 0	Blue	Blue
thread 1		
thread 2		
thread 3		
thread 4		
thread 5		
thread 6		
thread 7		

2 cycles

# Challenge 2 : Massive Thread Divergence

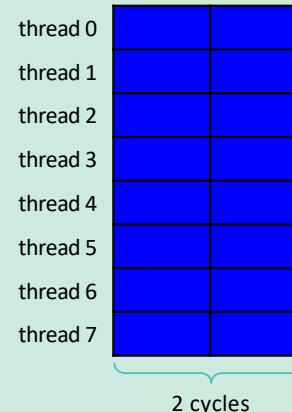
CS 0

```
if((threadIdx.x & 1) == 1){  
    code A;  
    if((threadIdx.x & 2) == 2){  
        code B;  
    } else{  
        code C;  
    }  
} else{  
    code D;  
    if((threadIdx.x & 2) == 2){  
        code E;  
    } else{  
        code F;  
    }  
}
```



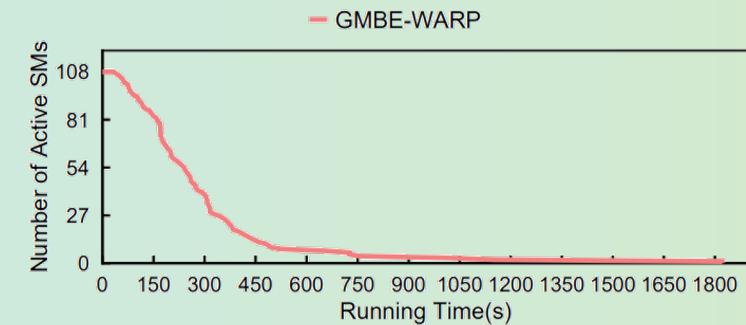
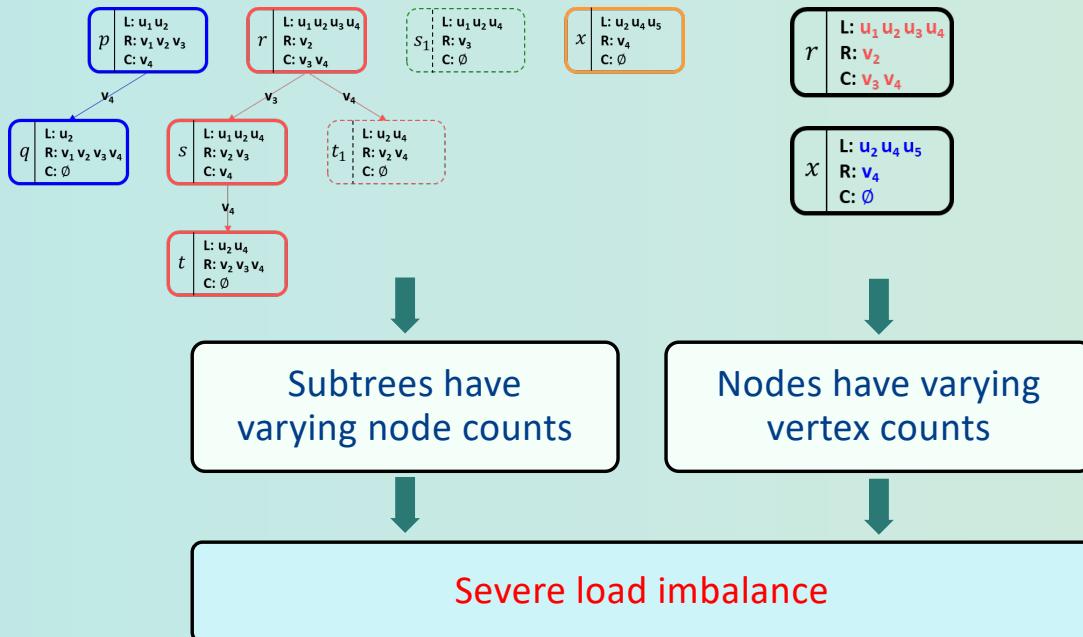
CS 1

```
code A;  
code A;
```



CS 0 requires more time because of the thread divergence.

# Challenge 3 : Load Imbalance



Active SMs rapidly decline over time in a straightforward algorithm.

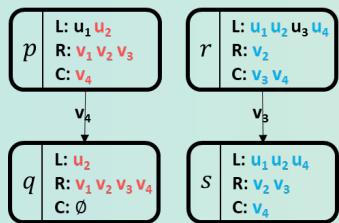
# GMBE

- Stack-based iteration with node reuse
- Pruning using local neighborhood sizes
- Load-aware task scheduling

# Idea 1 : Stack-based Iteration with Node Reuse

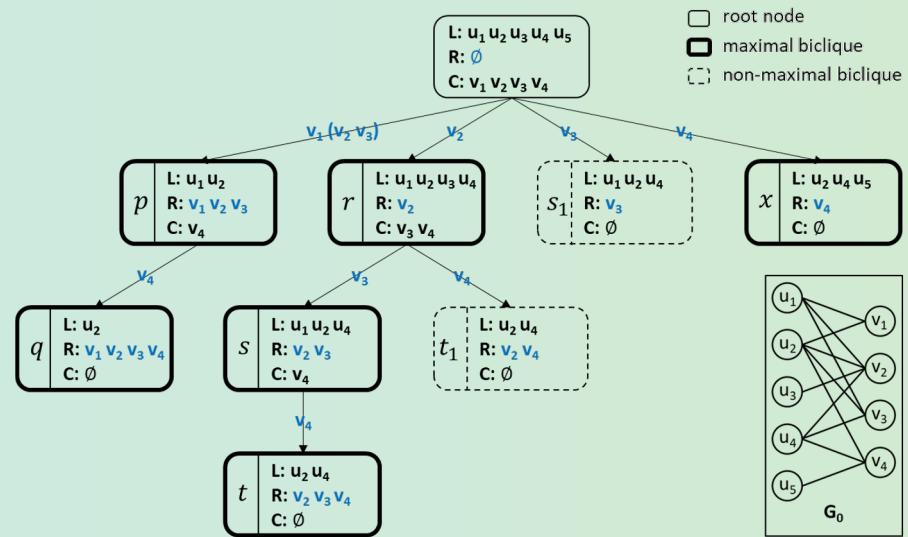
## ➤ Key Observation

- Vertices in the child node are always a **subset** of vertices in the parent node.



## ➤ Main idea

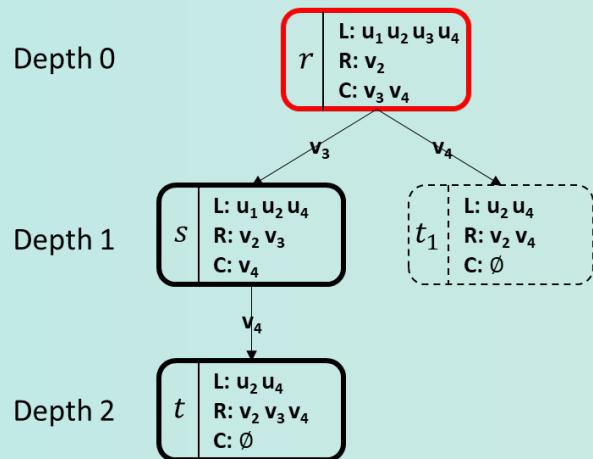
- We allocate memory for the root node and **reuse this memory** to derive all nodes within the subtree, resulting in a notable **reduction in memory usage**.



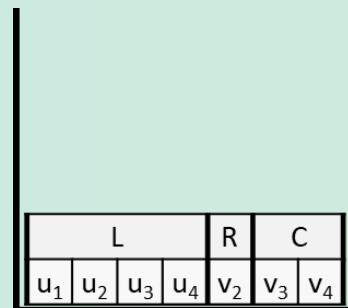
# Idea 1 : Stack-based Iteration with Node Reuse

## ➤ Step 1 : initialize node $r$

Enumeration tree rooted by node  $r$

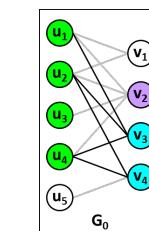


Memory usage for existing approach



Memory usage for GMBE

	$L_r$				$R_r$		$C_r$	
Vertex	$u_1$	$u_2$	$u_3$	$u_4$	$v_2$	$v_3$	$v_4$	
Depth	0	0	0	0	0	$\infty$	$\infty$	
Local neighborhood size	3				2			

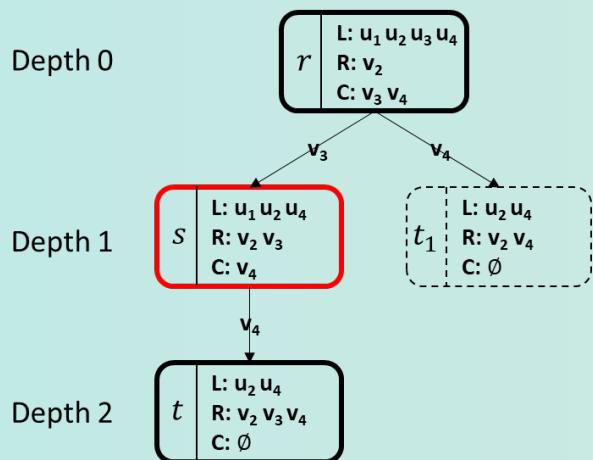


The local neighborhood size for  $v_3$  is 3 because  $v_3$  connects with  $u_1, u_2$ , and  $u_4$  in  $L_r$ .

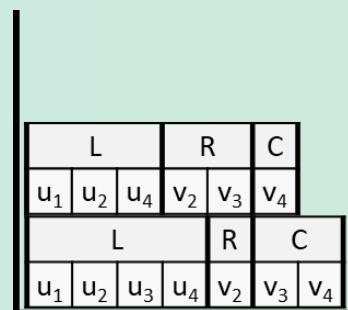
# Idea 1 : Stack-based Iteration with Node Reuse

## ➤ Step 2 : generate node s

Enumeration tree rooted by node  $r$



Memory usage for existing approach



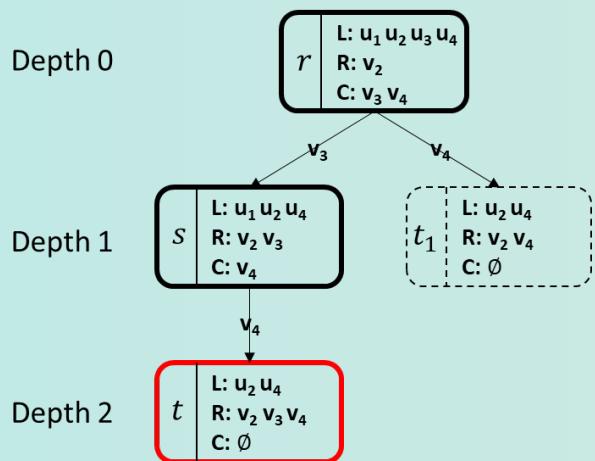
Memory usage for GMBE

	$L_r$				$R_r$	$C_r$	
Vertex	$u_1$	$u_2$	$u_3$	$u_4$	$v_2$	$v_3$	$v_4$
Depth	1	1	0	1	0	1	$\infty$
Local neighborhood size	3						2

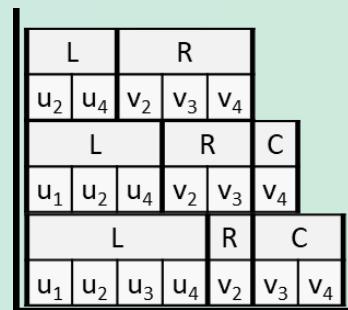
# Idea 1 : Stack-based Iteration with Node Reuse

## ➤ Step 3 : generate node $t$

Enumeration tree rooted by node  $r$



Memory usage for existing approach



Memory size is increasing  
↑

Memory usage for GMBE

Vertex	$L_r$				$R_r$		$C_r$	
Depth	1	2	0	2	0	1	2	
Local neighborhood size	2	2						

Memory size is fixed

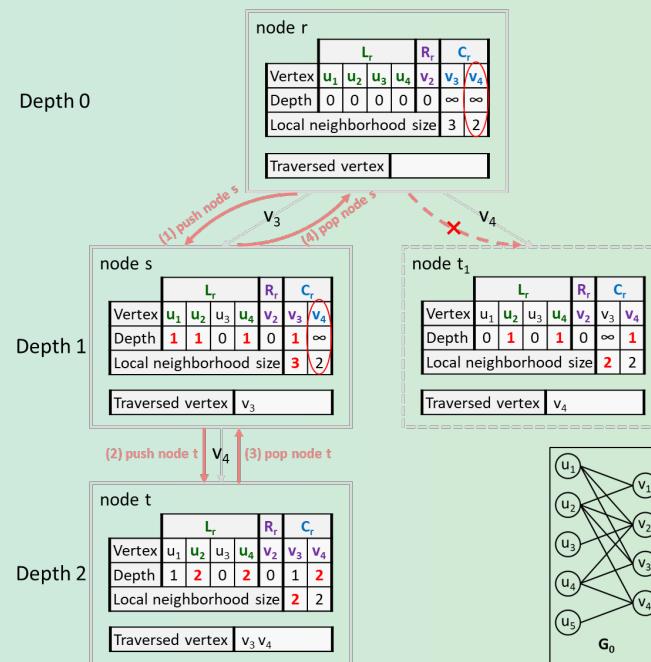
# Idea 2 : Pruning Using Local Neighborhood Sizes

## ➤ Key observation

- Local neighborhood sizes, as a **necessary intermediate result**, can be utilized for pruning.

## ➤ Main idea

- We prune useless candidates if **their local neighborhood sizes do not change after popping a traversed child node**. This approach **reduces thread divergence** by checking multiple local neighborhood sizes simultaneously..



GMBE proactively prunes node  $t_1$  by removing useless candidate vertex  $v_4$  at node  $r$  because the local neighborhood size (i.e., 2) for  $v_4$  does not change after popping node  $s$ .

# Idea 3: Load-aware Task Scheduling

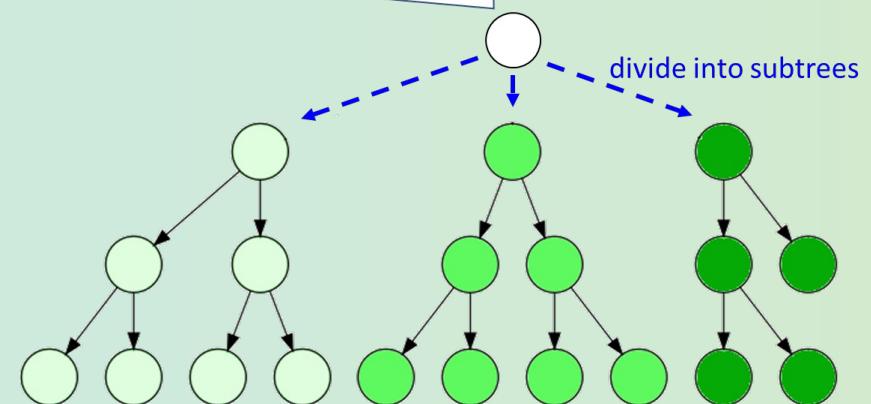
## ➤ Key observation

- Due to the imbalance of subtrees in MBE problems, directly mapping subtrees to computational resources leads to significant load imbalance.

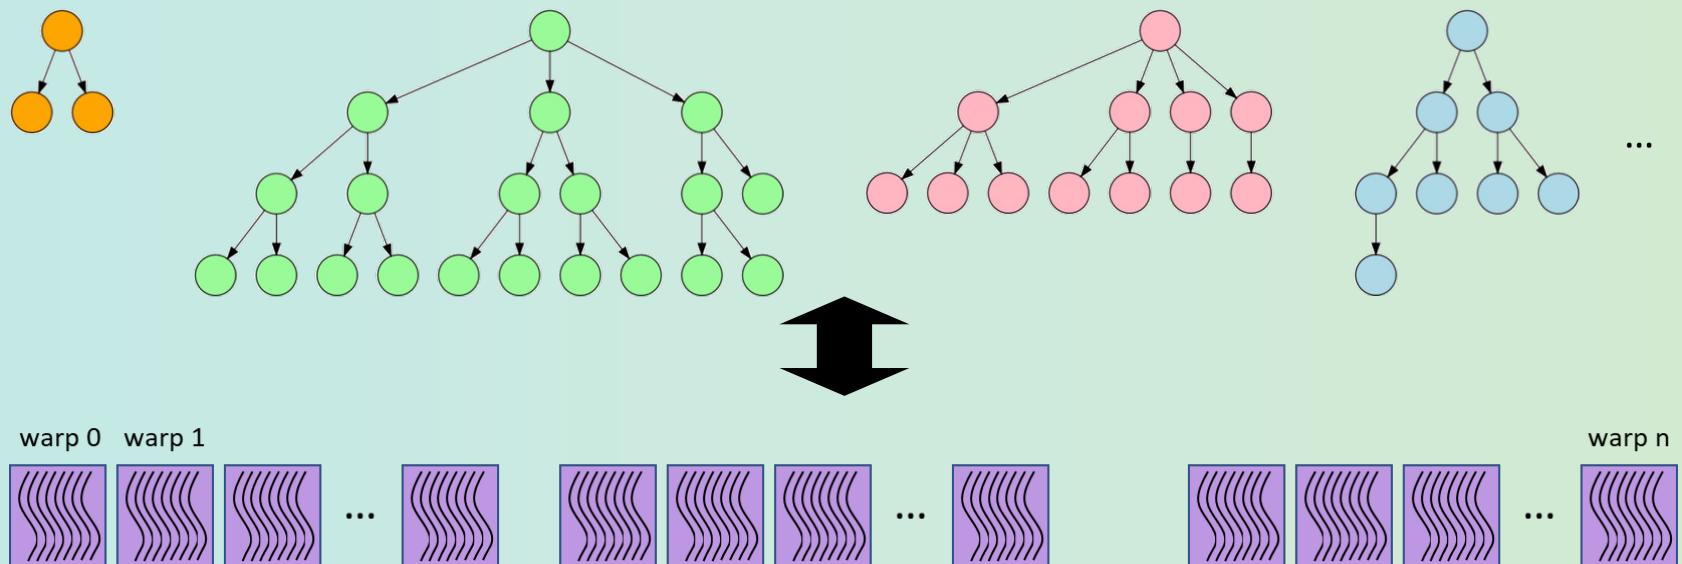
## ➤ Main idea

- Design two thresholds to **detect large subtree**.
- Dynamically **divide large trees** into multiple subtrees to **balance the workloads**.

For a subtree rooted at  $(L, R, C)$ . We will divide it if  $\min\{|L|, |C|\} > \tau_1$  and  $\min\{|L|, |C|\} \times C > \tau_2$

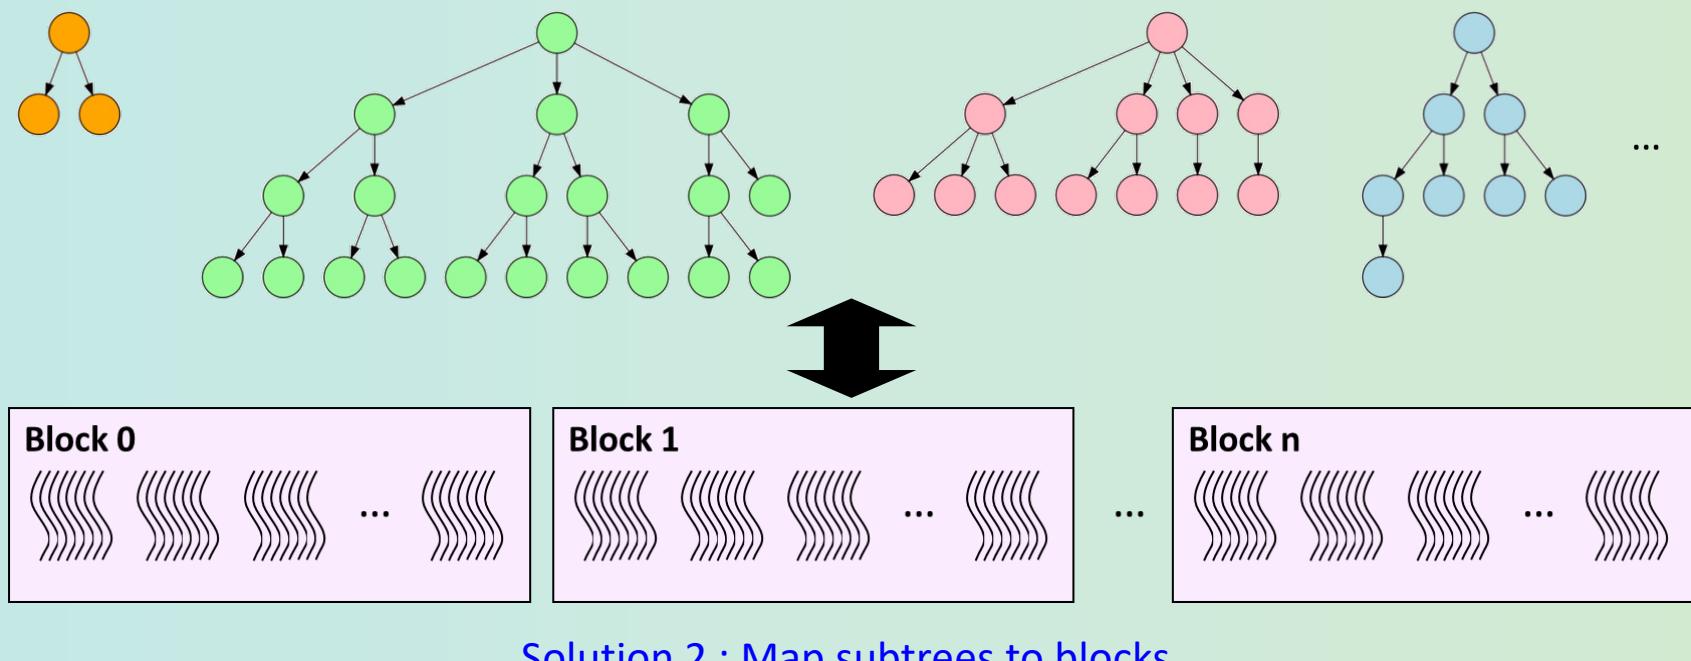


# Idea 3: Load-aware Task Scheduling

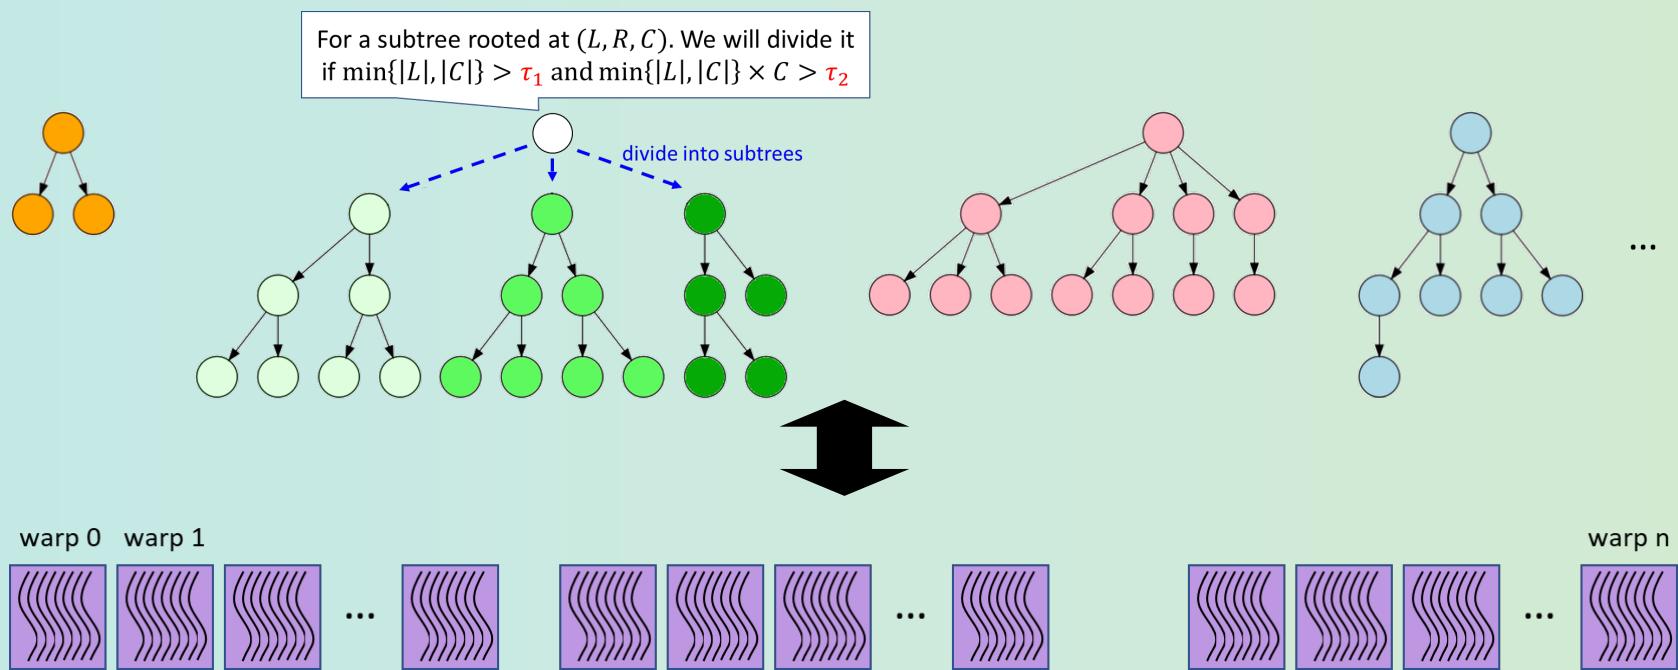


Solution 1 : Map subtrees to warps.

# Idea 3: Load-aware Task Scheduling



# Idea 3: Load-aware Task Scheduling



Solution 3 : Dynamically divide large tree into multiple subtrees and map subtrees to warps.

# Evaluation

- Overall evaluation
- Effect of optimizations
- Sensitivity Analysis

# Evaluation : Overall Evaluation

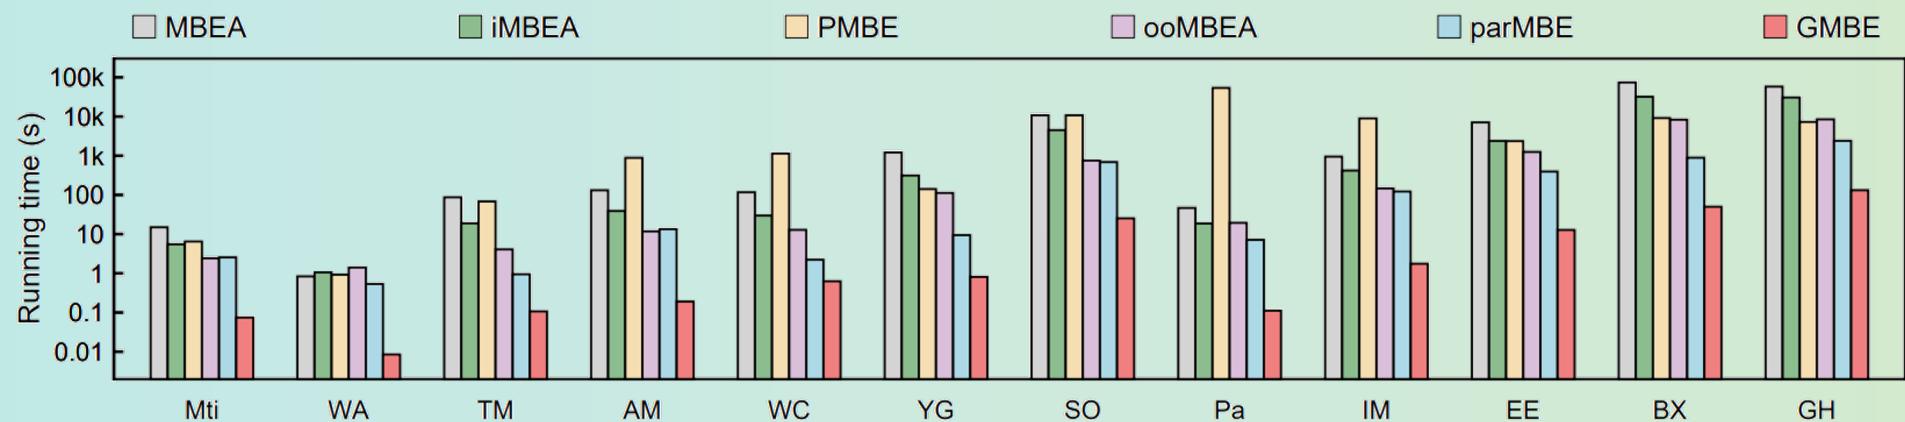


Figure 6: Overall evaluation (log scaled).

GMBE is  $3.5\times$ – $69.8\times$  faster than any next-best competitor on CPUs on all testing datasets.

# Evaluation : Effect of Optimizations

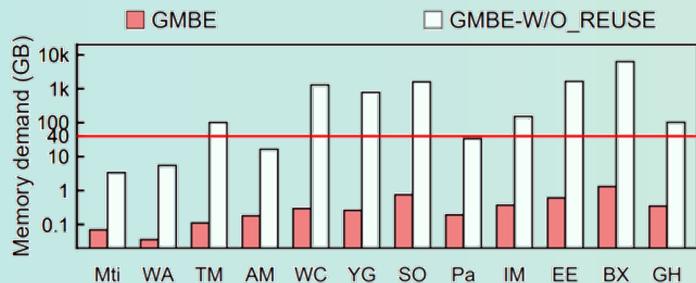


Figure 7: Effect of the node reuse approach (log scaled).

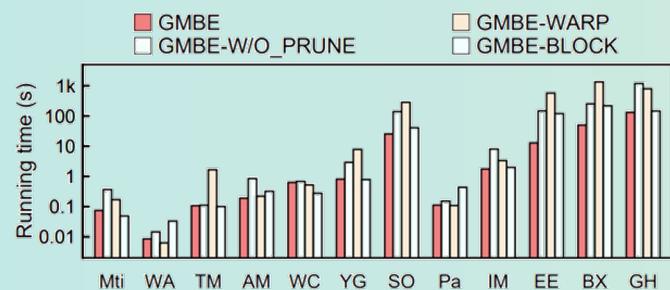
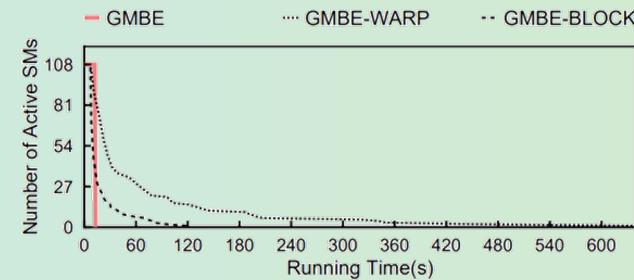
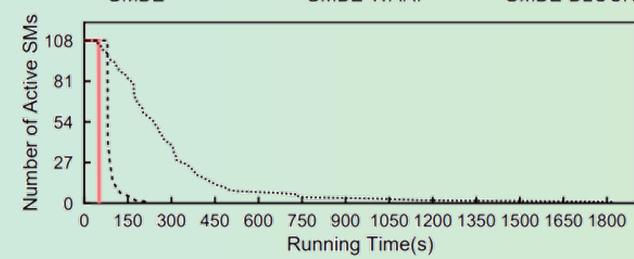


Figure 8: Effect of pruning approach and task scheduling approach (log scaled).

GMBE significantly reduces the memory usage, efficiently reduces the enumeration space, and successfully balance the workloads.



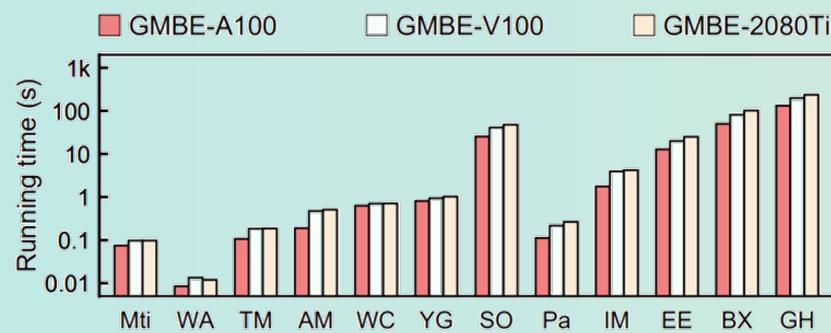
(a) EuAll



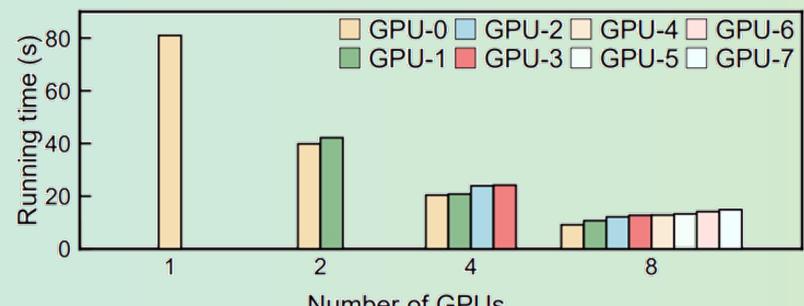
(b) BookCrossing

Figure 9: Comparison of runtime loads on SMs among GMBE, GMBE-WARP, and GMBE-BLOCK.

# Evaluation : Sensitivity Analysis



Adaptability on different GPU (log scaled).



(a) BookCrossing

Scalability of GMBE on a machine with multi-GPU.

# Q & A