

IMPRESS: An Importance-Informed Multi-Tier Prefix KV Storage System for Large Language Model Inference

Weijian Chen, Shuibing He, Haoyang Qu, Ruidong Zhang,
Siling Yang, Ping Chen, Yi Zheng [§], Baoxing Huai [§], Gang Chen




浙江大学
Zhejiang University



USENIX FAST 2025

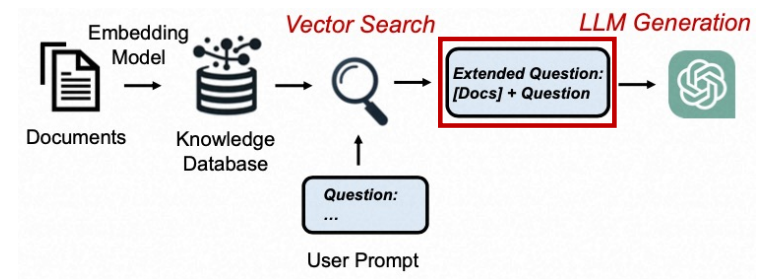
Large Language Model (LLM) Inference

- LLM has been applied in a range of fields





- context-rich **prefixes** + user **queries** = LLM **requests**
- Many **requests** **share identical prefixes**

Prompt:
[Instructions]
You are an AI chatbot. You are having a conversation with a human by following rules:
- You do not have a name.
- You are helpful, creative, clever, and friendly
...
[Examples]
Human: Hello, who are you?
AI: I am an AI chatbot. How can I help you?
...
[Question]
Human: Tell me about the second world war.




Prefix KV Storage System

- Shared prefix KVs can be **restored and reused**






* Assume a three-layer simple LLM

Time-to-First-Token (TTFT) can be reduced.



Prefix KV Storage System

- When shared prefix KVs needs to be stored into **SSD**



**Prefix KV loading from SSD to GPU
has become a new bottleneck**

Related Work

- Most existing systems store prefix KVs **only in GPU and/or CPU memory**

PromptCache-MLSys24, RAGCache-arxiv24, ChunkAttention-arxiv24, SGLang-arxiv23

Limited space in GPU and CPU memory
quickly becomes exhausted

- Pre-loads them into CPU memory based on the **scheduler's** predictions

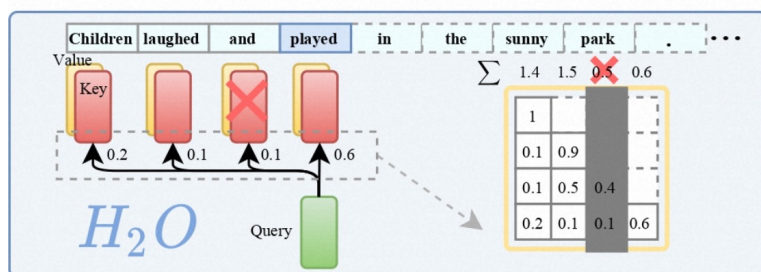
AttentionStore-ATC24

Limitations exists under high request
volumes or in preemptive scheduling

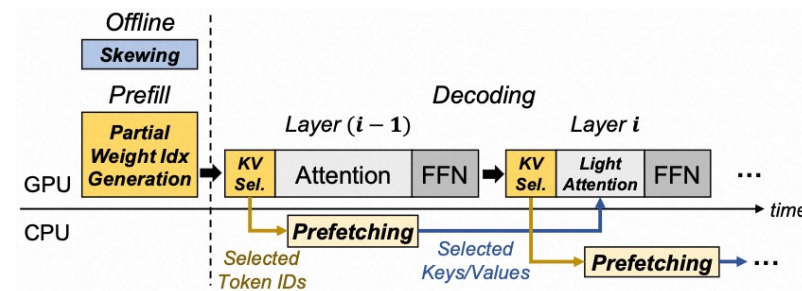
Is it possible to reduce KV data that needs to be loaded?

Opportunity from KV Importance

- Only preserve important KVs during decoding phase achieves the same level accuracy



H2O-NeurIPS23



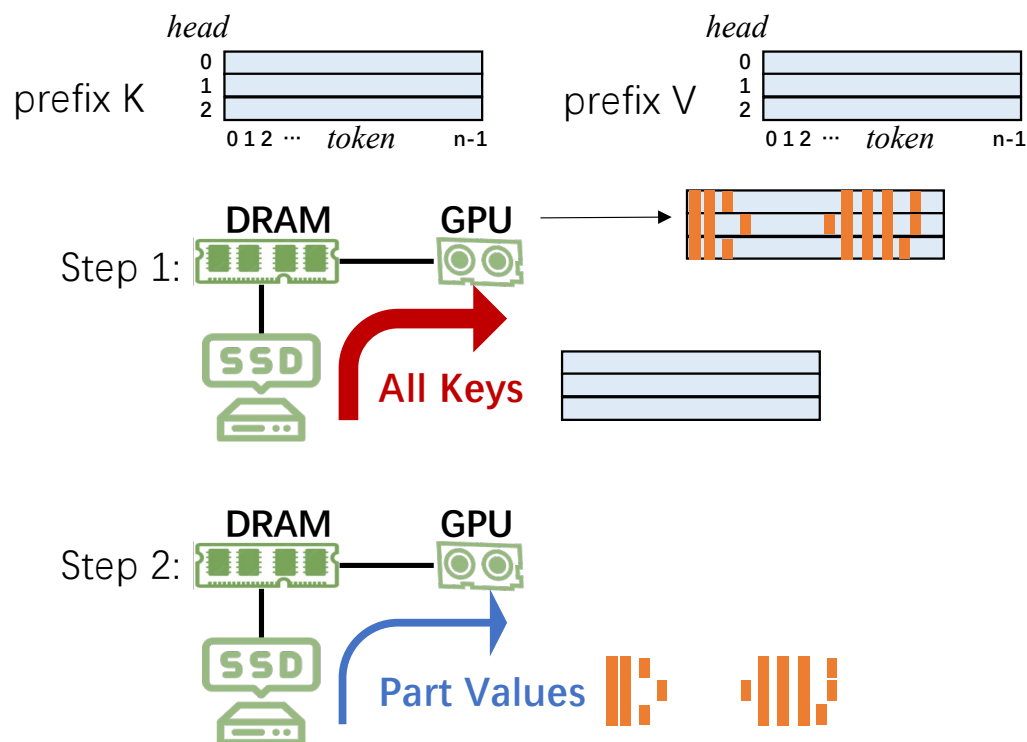
InfiniGen-OSDI24



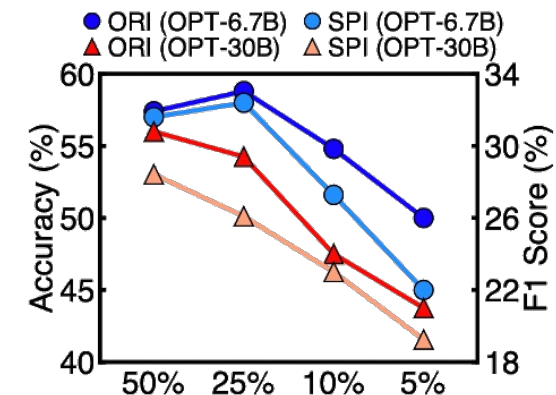
How about **only load** important KVs during **prefill** to reduce I/O bottleneck and TTFT?

Challenge 1

- A large amount of I/O is introduced to identify important KVs.



- Pre-determine important KVs?
Accuracy Drop.




- SPI: statically pre-determine importance
- ORI: original dynamically determine importance

Challenge 2

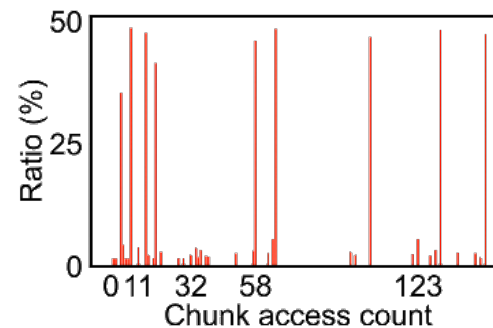
➤ The existing prefix KV **storage** and **caching** systems are **suboptimal** considering token's importance.

1. Storage: **read amplification**
(Each chunk contains a mix of important and unimportant KVs.)



(a) The ratio of important KVs within each chunk.

2. Caching: based solely on recency or frequency
(**ignore the importance** of KVs)

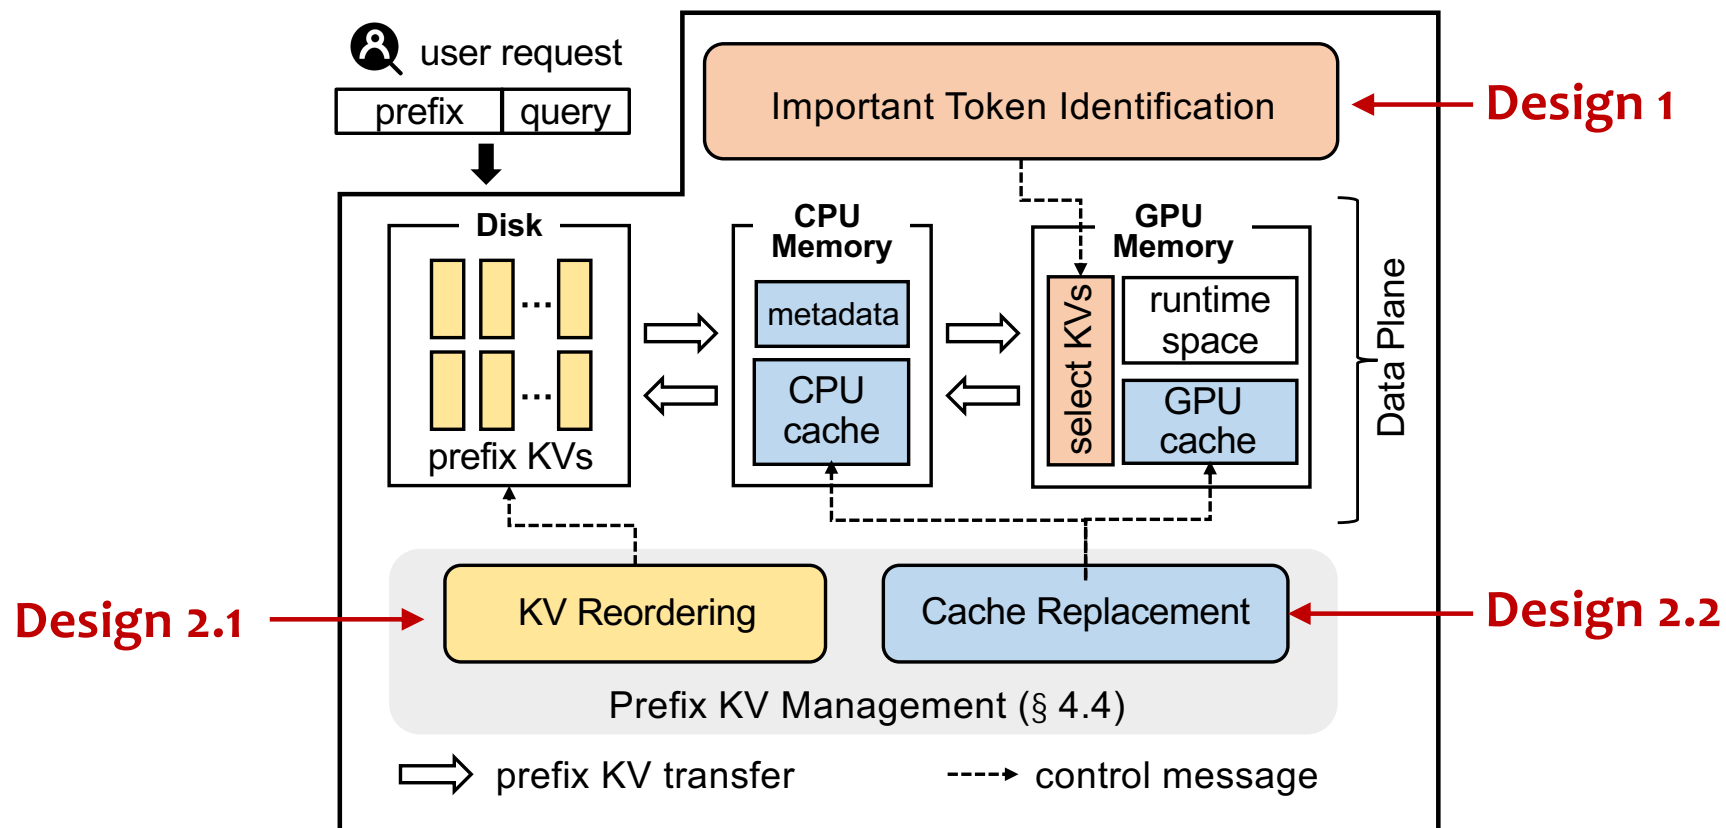


(b) Average ratio of important tokens in all chunks for a given chunk access frequency.

Outline


- Background & Motivation
- **Observation & Design of IMPRESS**
- Evaluation
- Summary & Conclusion

IMPRESS Architecture




Observation

- There is a high similarity in the set of important token indices across different heads within the same layer of an LLM.




Similarity measurement:

$$h_0 = \{0, 2\} \quad h_1 = \{0, 1\} \quad J(h_0, h_1) = \frac{|h_0 \cap h_1|}{|h_0 \cup h_1|} = \frac{1}{3}$$



(a) select the top 10% most important tokens

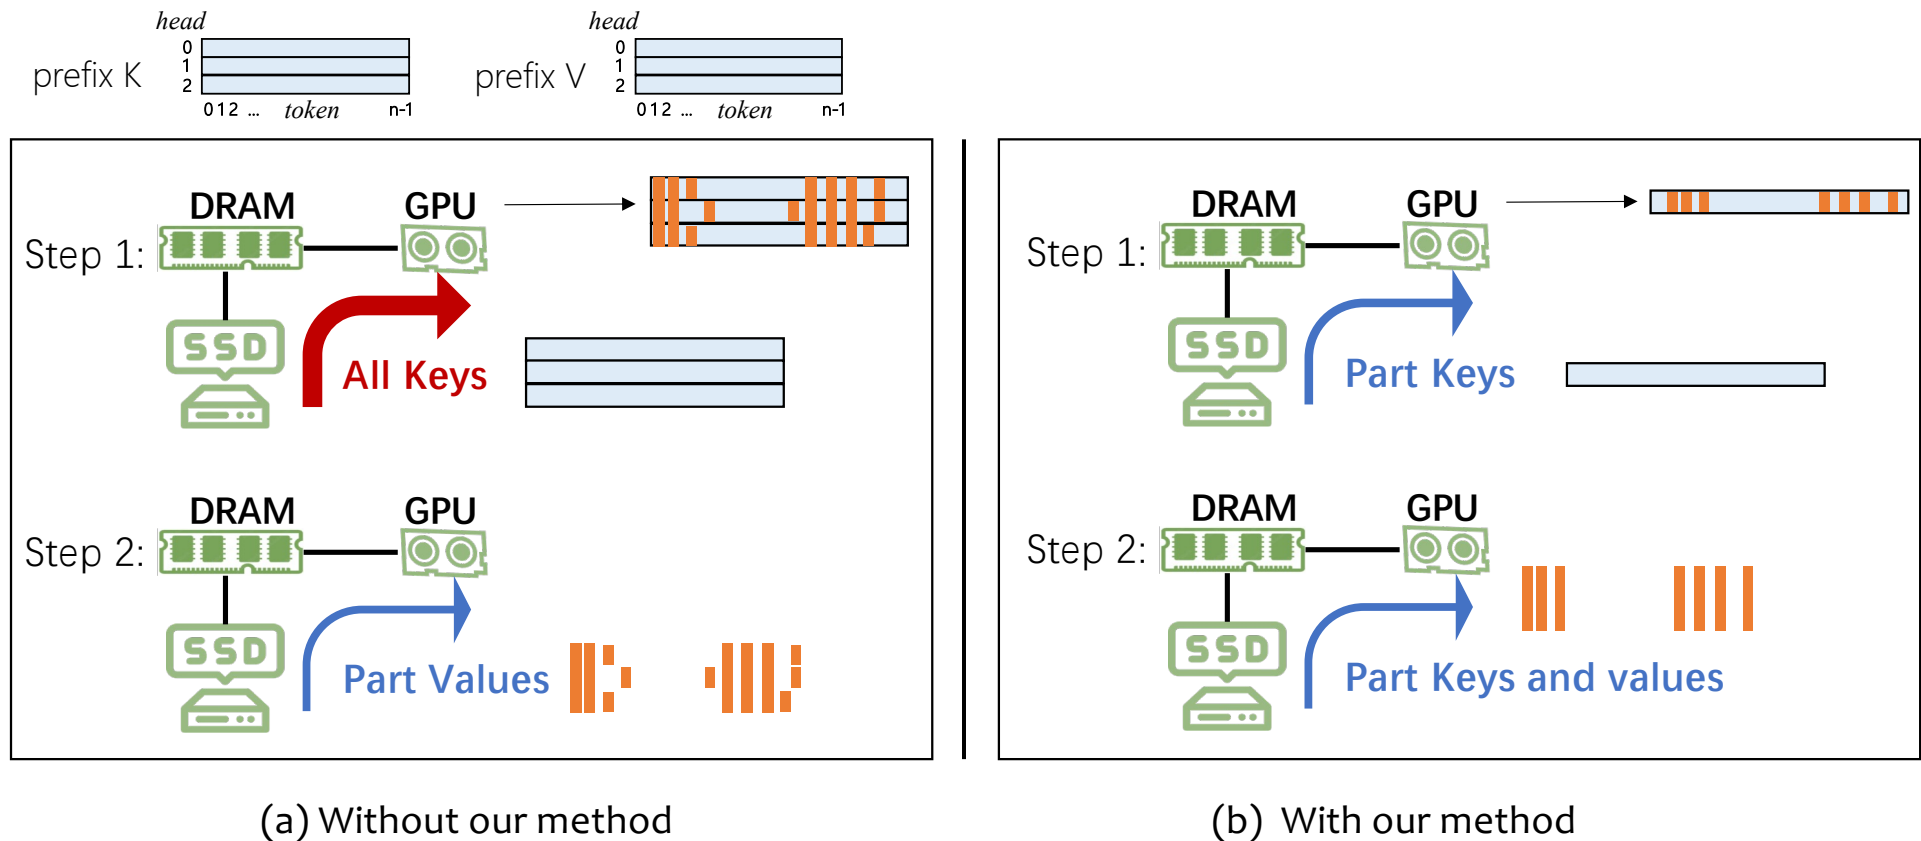


(b) select the top 40% most important tokens

The similarity of important tokens indices exists across different LLM scales and important KV ratios.

1 Similarity-Guided Important Token Identification



Key idea: Use the important token index set from a few selected heads to **approximate** the important token index sets for the remaining heads



2.1 KV Reordering

- Target: Reduce read amplification
- Key idea: reorder and repack important KV's into denser chunks

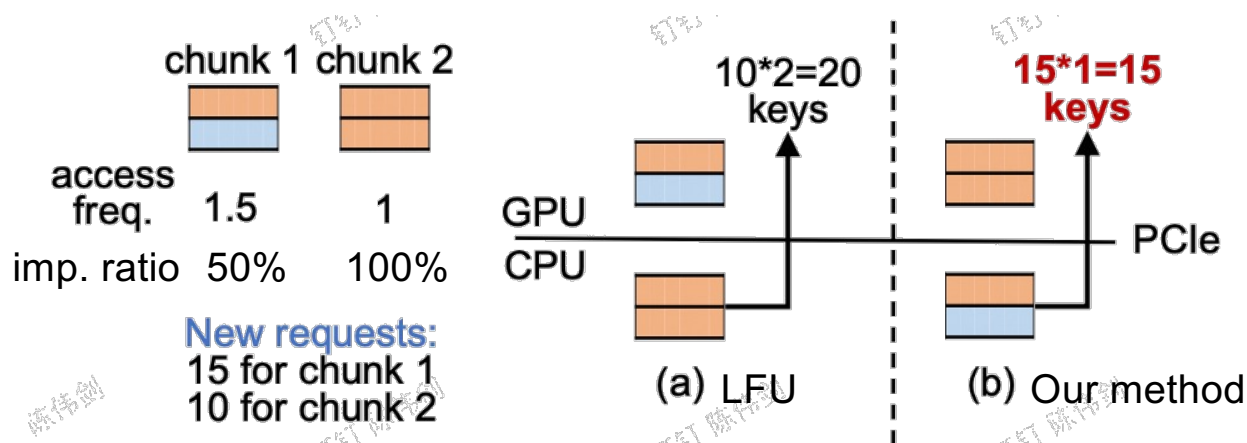
- Problem: KV reordering may destroy the radix tree structure by altering the token order
1. avoid cross-node reordering
 2. Add mapping list to recovery



2.2 Score-Based Cache Management

- Key idea: Data admission and cache replacement **based on scoring**.
- The score = the chunk access frequency * proportion of important KVs.

The higher the score, the higher the priority for admission into the faster medium cache.



score for chunk 1: $1.5 \times 50\% = 0.75$

score for chunk 2: $1 \times 100\% = 1$

Experimental Setup

➤ System configuration

CPU	2 × AMD EPYC 7763
GPU	1× NVIDIA A100 (80GB)
Memory & SSD	128 GB DRAM, 2TB SSD (5GB/s)

➤ Workloads and datasets

Datasets	PIQA, RTE, COPA, and OpenBookQA Prefix sizes: 55GB, 57GB, 64GB, 65 GB
Models	OPT-6.7B, OPT-13B, OPT-30B

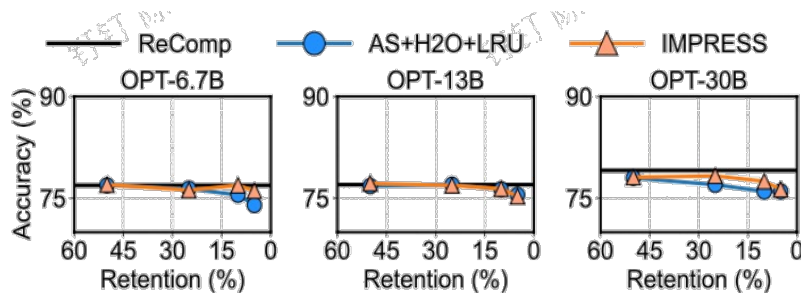
➤ Compared systems

ReComp	Recomputation without reusing prefix KVs
AS-like	AttentionStore with async KV loading, without scheduler
AS+H2O+LRU	Add H2O on top of AttentionStore with LRU
AS+H2O+LFU	Add H2O on top of AttentionStore with LFU
IMPRESS	Our three optimizations on top of H2O

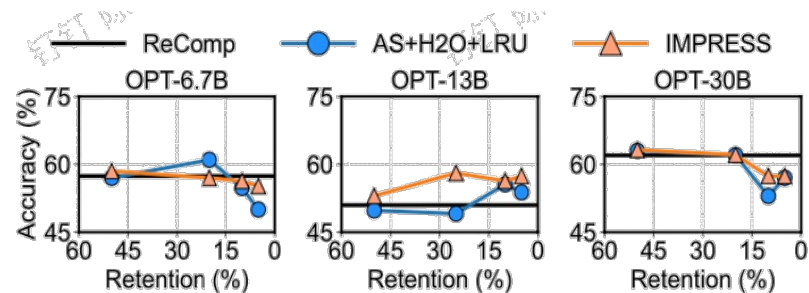
➤ Default settings.

- (1) cache size: 10GB GPU HBM, 32GB CPU DRAM
- (2) Chunk size: keys or values of 64 tokens


Model Inference Accuracy




(a) PIQA



(b) RTE



(c) COPA

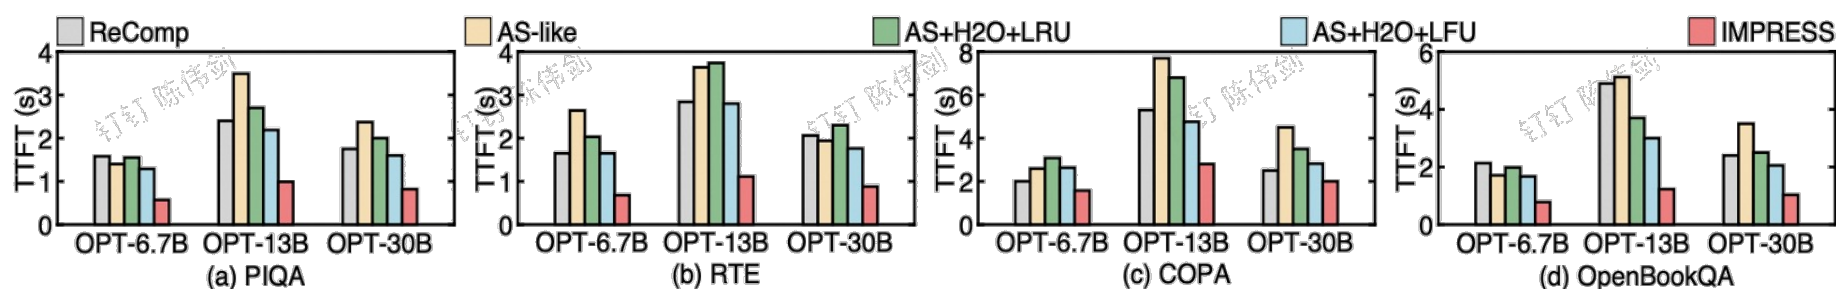


(d) OpenBookQA


Compared to ReComp, the average inference accuracy drop is less than 0.2%

Time-to-first-token (TTFT)


✓ TTFT



✓ I/O time




✓ Individual technique




IMPRESS outperforms alternatives, with a $1.2\times\sim 2.8\times$ improvement over SOTA solutions, due to a $1.5\times\sim 3.8\times$ reduction in I/O time.


Sensitivity Analysis



Various chunk sizes



Various dataset scales



Results on Llama models

IMPRESS outperforms the leading alternative on various cases.

More evaluations: checkout our paper

Summary & Conclusion

➤ Problem

- I/O becomes the bottleneck when shared prefix KVs are loaded from SSD for LLM

➤ Key idea

- Only load important KVs during prefill phase

➤ Challenges

- A large amount of I/O is introduced to identify important KVs
- Storage and caching systems are suboptimal

➤ Techniques in iCache

- Similarity-Guided Important Token Identification
- KV Reordering & Score-Based Cache Management

➤ Results

- IMPRESS outperforms the alternatives with the same level of inference accuracy

Thanks & QA

IMPRESS: An Importance-Informed Multi-Tier Prefix KV Storage System for Large Language Model Inference



浙江大学
Zhejiang University



HUAWEI

HUAWEI CLOUD



Contact Email: weijianchen@zju.edu.cn



ISCS Lab: <https://shuibing9420.github.io>